

The Sources of Researcher Variation in Economics

Nick Huntington-Klein*

Claus Pörtner

We have 146 research teams of economists independently use the same data source to answer the same question about the causal effect of a policy. Each team performs the task three times, first allowing free choice of how to answer the question, second requiring all researchers to use a shared research design, and third providing pre-cleaned data. We find considerable variation across researchers in reported sample sizes, sample definitions, and modeling choices, although actual estimated effects show muted variation. Levels of researcher variation were most heavily influenced by data preparation and cleaning, with much smaller roles for researcher background, modeling and estimation decisions, and peer review. This implies that data preparation and cleaning should receive considerably more attention in researcher training and the evaluation of research.

1 Introduction

The social and behavioral sciences produce a staggeringly large flow of empirical results. A responsible reader of this literature should wonder how much they can trust a given study, given the potential for errors, fluke results, or intentional manipulation. Any responsible producer of this literature, however, is well aware that even a researcher doing their best to avoid these problems must make hundreds of choices in the process of collecting and cleaning data, planning their estimation, and coding their analysis, in other words “researcher degrees

*Corresponding author Nick Huntington-Klein, nhuntington-klein@seattleu.edu, +1 (206) 296-5815. Department of Economics, Seattle University, 901 12th ave., Seattle, WA, 98122. Huntington-Klein and Pörtner are the project organizers. This project was supported by the Alfred P. Sloan foundation grant G-2022-19377. The Seattle University IRB determined this study to be exempt from IRB review in accordance with federal regulation criteria. Many thanks to Kian Farzaneh, Amrapali Samanta, and Erica Long for research assistance, to the researchers Mira Chaskes, Jennifer A. Heissel, Elaine L. Hill, Rajius Idzalika, Joshua D. Merfeld, and Ethan Sawyer, who contributed but did not want an authorship slot, to the researchers who wished to remain anonymous, and to the researchers who enlisted in the study but were not eligible or were not able to complete all three rounds of the project.

of freedom” (Simmons, Nelson, and Simonsohn 2011). Even if Researcher A’s choices would stand up to scrutiny and argument from a reviewer or reader, Researcher B with the same goal, data, and skills might have reasonably chosen in a different way that would also stand up to scrutiny. If A and B’s choices lead to different results, but only one of them performs the study, then this is a source of largely arbitrary variation in the collection of published results. Estimates in one context suggest this variation might outweigh the population variation we typically consider when estimating standard errors (Holzmeister et al. 2023).

This study looks at the impact of researcher degrees of freedom on results, and also attempts to isolate researcher degrees of freedom at different stages of the research process to try to determine where researcher choice is most varied, and most strongly influences results. We do this using a “many-analysts” design where multiple researchers attempt the same research task, specifically a task that is common across applied econometrics: estimating the causal effect of a policy that is implemented at a specific period of time and affects some people but not others. We look at differences between researchers in the results as well as the analytic and data cleaning choices made.

We expand on a typical many-analysts design by introducing multiple iterations of analysis, each time restricting the amount of choice that researchers can make and so reducing researcher degrees of freedom. This allows us to observe the overall amount of variation in estimates between researchers, as is common in many-analysts designs, and also to separately evaluate the influence of choice in research design and in data cleaning, and the impact of peer review.

We find meaningful differences in the ways that different researchers approach the same task. Some of these differences come from decisions that would receive scrutiny from a reader, like the research design and choice of control variables. Other differences came from sources where researchers did choose differently but a reader might not recognize that a consequential decision had been made, like in the functional form of the control variables or on a number of data cleaning or sample limitation decisions. When researchers were forced to all use the same research design, results became more similar, especially among the researchers most familiar with the subfield of the research task. Researcher agreement increased sharply when pre-cleaned data was provided to researchers, implying that data cleaning decisions are a major source of variation between researchers. Development of more mature and standardized data cleaning procedures, and increased visibility for data cleaning, may have a meaningful impact on the consistency and believability of results in applied microeconomics.

1.1 Previous Work on Research Reliability and Researcher Degrees of Freedom

In economics, suspicion about empirical results is not new (Leamer 1983). The most recent wave of concern is inspired by discussions, originating in the field of psychology, of the “replication crisis.” Studies on the replication crisis show that a high percentage of studies cannot be replicated when tested using new data Camerer et al. (2016), that study code and data is not available or does not reproduce the published results (Herbert et al. 2021), or that

“policing replications” that test sensitivity of published results are rare (Ankel-Peters, Fiala, and Neubauer 2023).

However, this greater body of replication work takes an existing study as a baseline and asks whether it is robust to re-evaluation in some way. Questions about researcher degrees of freedom are not about whether a given study can be challenged, but whether a different researcher performing the same study would have done it differently if they had been the person to perform it, without taking the original as a baseline.¹ There is some research on this topic in regards to researcher identity or political orientation, as in Jelveh, Kogut, and Naidu (2024), or personality characteristics, as in Sulik et al. (2023). But a common way to empirically study researcher degrees of freedom is using a many-analysts design.

The many-analysts design,² popularized by Silberzahn et al. (2018), gives the same data set to multiple teams of researchers and has them independently try to answer the same research question.

Many-analysts studies have now been carried out in many fields, including microeconomics (Huntington-Klein et al. 2021), finance (Menkveld et al. 2021), religion (Hoogeveen et al. 2023), neuroimaging (Botvinik-Nezer et al. 2020), political science (Breznau et al. 2021), machine learning (Chen and Cummings 2024), ecology and evolutionary biology (Gould et al. 2023), psychology (Boehm et al. 2018; Bastiaansen et al. 2020; Schweinsberg et al. 2021), and medical informatics (Ostropolets et al. 2023), among others.

With little exception, many-analysts studies find that there *is* meaningful variation in both methods and conclusions across researchers. Holzmeister et al. (2023) finds that researcher variation in design and analysis likely outweighs population variation in effects.

These studies vary considerably, however, in the extent to which they can establish the source of that researcher variation or suggest policies that might reduce it. Establishing that there is variation is important, but is of limited impact if we do not understand why it is there or what we can do about it. Further, many-analysts results may not even imply a problem if not carefully performed. Variation in the original Silberzahn et al. (2018) study may be largely explained by the research question not being made sufficiently clear to researchers (Auspurg and Brüderl 2021), and skipping standard meta-analytic practice may overstate variation between researchers by being too sensitive to outlier estimates (Auspurg and Brüderl 2023).

The ability to explain variation between researchers, rather than just show that variation exists, is limited by the size of these many-analyst studies. As will be explored in Section Section 3.2.2,

¹These two fields intersect, and some failures to replicate in replication studies may be due to researcher degrees of freedom, where both the original study and the replication made reasonable choices but found different results (Bryan, Yeager, and O’Brien 2019). The difference in framing here is that the replication literature views the differing choices as a challenge to the validity of the original results, while the researcher degrees of freedom framing views both as part of a universe of reasonable results, assuming both analyses are defensible.

²Many-analysts designs are sometimes referred to as “crowdsourced” science.

this study pursued a sample of at least 90 researchers so as to have acceptable power to explain differences in variation. Pérignon et al. (2022), in looking at the sources of reproducibility variation using many teams, used a design with 1,000 tests to replicate in order to adequately power comparisons. Since participation in a many-analysts study takes considerable time and effort, sample sizes are often well below even the aforementioned 90, which may explain why many studies do not attempt to decompose the variation in effects they find between sources. These smaller sample sizes can produce acceptable statistical power for some tests but not others, and explaining variation or agreement in effects between researchers generally demands a larger sample than showing the existence of meaningful variation or showing a difference in rates of making a particular research decision. Many-analyst studies that aim to explain the sources of variation between researchers either do so despite the low-power issue, gather larger samples of researchers, or select analyses that produce adequate power despite small samples.

Among studies that attempt to explain researcher variation, there are three common explanations explored. The first of these is in the difficulty of the research task, with some studies showing less researcher agreement in more complex or difficult-to-analyze scenarios (Menkveld et al. 2021; Ortloff et al. 2023). A second source is researcher experience or characteristics. Menkveld et al. (2021) find that higher-quality teams (with more experience, seniority, publishing success, and/or people) agreed more. Ortloff et al. (2023) find that experienced researchers tended to draw more abstract codebooks and conclusions than students, and Broderick, Giordano, and Meager (2020) find that replicators with more coding skill found more errors in original work. Breznau et al. (2021), however, found that researcher characteristics explained only a small share of the variation in results.

A third factor used to explain variation is peer review or evaluation. Seeing the actual impact of review requires that researchers be able to revise their work after receiving it, as in Menkveld et al. (2021), who find that review increases agreement. In some cases there is no chance to revise so we cannot see the impact of peer review, but instead outside evaluation is used as a measure of researcher quality. In this vein, Gould et al. (2023) find that peer review scores do not predict whether a given researcher produces an outlier result.

Outside of many-analyst designs, there are studies that use simulation to try many combinations of analytical or data-cleaning choices and examine the resulting variation in estimates. This approach is similar to a many-analysts design in that they look at variability in potential research choices and, often, try to explain variation in effects estimates using those choices. They differ in that they are necessarily limited to the set of research decisions that the project organizers consider ahead of time (which constrains the universe of possible decisions but also makes interpretation of the results far more clear), and typically consider all combinations of decisions equally, rather than favoring combinations an actual researcher would choose. Of these studies, the closest to the present study is Klau et al. (2023), which evaluates the sensitivity of results in an observational psychological data set to different data preprocessing and modeling choices. They try multiple combinations of reasonable preprocessing and modeling choices, using simulation to iterate through the universe of potential choices, and find

significant variation in effects over reasonable preprocessing and modeling choices. A similar attempt to separate researcher variation into modeling and preprocessing components is also done in a many-analysts design in Huntington-Klein et al. (2021), although in a limited way.

This study’s design attempts to evaluate multiple of these sources of variation using a staged design, similar to Pérignon et al. (2022). The different stages allow different levels of researcher choice along the lines of interpretation of the research question, research design, and data preparation, as well as randomized peer review incorporating these mechanisms proposed by the literature and responding to the critique of Auspurg and Brüderl (2021). Researcher characteristics are collected, as well, allowing for exploration of the researcher-characteristics source of researcher variation, although not in a controlled way. We do not address the difficulty of the research task as a potential source of researcher variation in this study.

2 Design

In this study, we attempt to isolate the influence of several different potential sources of researcher variation by having the same set of researchers complete the same research task at least three times. We refer to these main research tasks as Task 1, Task 2, and Task 3. Following each task there is also a round of peer review and an opportunity to revise work.

Task 1 gives each researcher a large amount of freedom in terms of how they plan to complete the research task. Each successive task removes a degree of freedom from the researcher and further specifies how the analysis is to be performed. The intuition behind this design is that if the removal of a specific kind of researcher freedom meaningfully reduces the variation in results between researches, then that degree of freedom is a meaningful contributor to researcher variation.

The following goals and instructions are shared across all tasks:

- Estimate the causal effect of a policy on a specified outcome, among the group affected by that policy (see Section 3.1 below for more details).
- Use American Community Survey (ACS) data to estimate the effect, using data no older than 2006 and no newer than 2016.
- Procure ACS data from IPUMS (Ruggles et al. 2024), selecting only one-year files and using harmonized variables.
- Optionally, combine the ACS data with a data set on the presence or absence of other relevant policies, provided by the organizers.
- Use a statistics package or language that allows results to be immediately replicated.

Researchers were also given background information on the policy itself and its eligibility criteria, guidance on how to use the IPUMS website, instructed to use assistants for any work they would normally use assistants for, and to complete their analysis as though it had been their own idea, rather than attempting to match or not-match other researchers, or asking the project organizers how they would like the analysis to be performed.

These instructions comprise the entirety of the limitations on researchers in Task 1. Tasks 2 and 3 specified the task further and removed researcher degrees of freedom.

- Task 2 specified the research design more precisely. Instead of allowing any research design to identify the causal effect of interest, Task 2 gave specific definitions for which individuals comprised a “treated” group and which comprised an “untreated” comparison group.³ Then, it instructed researchers to estimate the effect by comparing how outcomes for the “treated” group changed from before policy implementation to afterwards against how outcome for the “untreated” group changed. This can be thought of as a difference-in-differences style design, although the phrase “difference-in-differences” was not used in the instructions.
- Task 3 uses the same research design limitations of Task 2, but also provides a pre-cleaned data set, prepared by the organizers. The data set offered a pre-prepared treated/untreated-group indicator as specified in Task 2, limited the data set only to the treated and untreated group, prepared and cleaned all variables in the data set that did not already come pre-cleaned, handled missing-data flags, merged in state policy data, and offered standardized simplified recodings of demographic variables. Researchers were instructed to not further clean the data or limit the sample.

Comparison of the researcher output between Task 1 and Task 2 is intended to show the researcher variation introduced by either an imprecise statement of the research question, as in Auspurg and Brüderl (2021), or due to differences in research design choices.

Comparison of the researcher output between Task 2 and Task 3 is intended to show the researcher variation introduced by decisions made in the data cleaning and variable definition process. A researcher following the Task 2 instructions should arrive at the same sample size, number of treated individuals, and number of untreated individuals as in Task 3, as well as the same definition for the outcome variable.⁴ Differences in the data set and in the results between Task 2 and Task 3 should be a result of differences in the data cleaning and preparation process.

³Although eligibility criteria for the policy were explicitly given in Task 1, Task 2 further limits the treated group by narrowing the acceptable age range. The limitation was more impactful for defining the untreated comparison group, though. Many researchers did use a treated/untreated group approach in Task 1 before it was specified in Task 2, but different researchers defined the untreated group in highly diverse ways, as will be shown in the Results section.

⁴The Task 2 instructions do leave some leeway for definition of some variables, in particular control variables like education or race, which have a specific recoded version available in Task 3 that are not specified in the Task 2 instructions. However, the definitions of the treated and untreated comparison groups should be the same between Task 2 and Task 3.

Following each of the research tasks, researchers engage in a round of peer review. 2/3 of researchers are randomly assigned to peer review, and 1/3 do not engage in peer review. Those in peer review are randomly assigned in pairs. Those pairs performed a blind review of each others' work, and provided a written assessment of that work. Reviewers were instructed to produce a review "as though (they) were the reviewer of a journal article," and to judge the work as though they were reviewing for a journal where a study of this kind "could be published if the work was of high quality."

Following peer review, all researchers have an opportunity to revise their work in light of the peer review (or for any other reason). Importantly, revision is not mandatory, nor is satisfying one's peer reviewer, and the majority of researchers did not choose to submit revisions.

Notably, this form of peer review does not exactly match what is typically done in peer review work for journal publications. In particular, revision is non-mandatory, all reviewers have themselves completed a study with the same goal and data and so have extensive background information, and all reviewers are themselves also reviewed by the same person. These features will all affect interpretation of the peer review results. In particular, the non-mandatory nature of the peer review means that the between-round revision work is only visible for a small subset of the researchers, and the paired nature of the reviews means we cannot separate the effect of being reviewed from the effect of reviewing someone else.

Following each research task and revision, researchers filled out a survey about their work.⁵ This survey asked them to report their findings, additional information like sample size and standard errors, and choices made in the process of doing the analysis like sample restrictions, treated-group definitions, estimator, and standard error adjustments. Researchers were also asked to justify why they had made these choices.

This research design and analysis plan has been preregistered (Portner and Huntington-Klein 2022). Analyses that were not preregistered will be noted in the results section as they are performed. Full instructions for each task, as well as post-task survey text and the peer-reviewing instructions, are available in the online appendix.

3 Data

3.1 The Focal Research Task

In all research tasks, the specific goal given to researchers was:⁶

⁵Note that the design of this study, and this survey, predates Sarafoglou et al. (2024) and so does not follow it.

⁶Full instructions are available in the online appendix.

Among ethnically Hispanic-Mexican Mexican-born people living in the United States, what was the causal impact of eligibility for the Deferred Action for Childhood Arrivals (DACA) program (treatment) on the probability that the eligible person is employed full-time (outcome), defined as usually working 35 hours per week or more?

DACA was implemented in 2012. Examine the effects on full-time employment in the years 2013-2016.

In simple terms, this asks researchers to estimate the impact of the DACA program on the probability that those eligible for the program usually work 35 hours per week or more in the years 2013-2016.⁷

Researchers, many of whom are not from the United States and so may not be familiar with DACA, are given further background information about the DACA program:

- DACA allowed undocumented immigrants who were accepted into the program to have legal work authorization for two years without fear of deportation, and also allowed them to apply for drivers' licenses or other forms of identification. People could reapply after the two years expired, and many did.
- Applications for the program opened on August 15, 2012, and over the first four years of the program's existence, over 900,000 applications were received, about 90% of which were approved.(U.S. Citizenship and Immigration Services 2016)
- While the program was not specific to immigrants from any origin country, because of the structure of undocumented immigration to the United States, the great majority of eligible people were from Mexico.

Researchers were also given information on the eligibility criteria for DACA, which was intended to apply only to a specific subset of undocumented immigrants who arrived in the United States as children, and not to all undocumented immigrants. Eligible people must:

- Have arrived in the United States before their 16th birthday.
- Not have had their 31st birthday as of June 15, 2012.
- Have lived continuously in the United States since June 15, 2007.
- Were present in the United States on June 15, 2012 and did not yet have legal status (either citizenship or legal residency) during that time.

⁷There are several existing papers that use the same ACS data set to identify the effect of DACA on various outcomes. The design used in Tasks 2 and 3 was most directly inspired by Amuedo-Dorantes and Antman (2016), although the designs do not match exactly, and the outcomes of interest are not the same. Researchers are informed that such previous studies exist and that they can optionally look into previous studies for background as they would normally do when performing research, although no specific previous study is listed. The instructions emphasize that any previous study does not constitute a "right answer" that researchers should be trying to match.

An additional eligibility requirement was mistakenly omitted from the Task 1 instructions, but was included for Tasks 2 and 3:

- Eligible people must have completed at least high school (12th grade) or be a veteran of the military.

In addition to this information about the policy itself and the effect that researchers are supposed to identify, researchers were also given instructions about the data set to use and how to procure it, as well as some details on usage of the data:

- Data should come from the American Community Survey (ACS), using data no older than 2006, and no newer than 2016.
- In addition, a file of state/year-level data was provided including labor market data and the presence or absence of different immigration policies in different years. Immigration policy data comes from Urban Institute (2022).⁸

ACS data should be procured from the IPUMS website (Ruggles et al. 2024), specifically selecting one-year ACS files and harmonized variables. Written and video instructions were included showing how to select data samples and variables on the IPUMS website.

- Researchers were not told which specific variables to use to determine eligibility status, but they were given guidance onto how to find relevant variables (like looking at the Person → Race, Ethnicity, and Nativity page to find variables relevant to ethnicity, birthplace, citizenship, and year of immigration).
- Several relevant features of the ACS that may affect analysis were emphasized: (a) ACS is a repeated cross-section, not a year-to-year panel data set, and (b) ACS does not list the month that data was collected in, so it is not possible to distinguish whether a given observation in 2012 is from before or after the policy was implemented, and (c) we do not actually observe in ACS whether a given person is enrolled in DACA, so we assume that all eligible people who are ethnically Mexican and Mexican-born are treated.

Finally, researchers were instructed to keep track of any variables used to limit their sample download on IPUMS, and to review the survey where they would be reporting their results before beginning their analysis.

From there, researchers were given free reign to complete the analysis as they thought most appropriate, including their own choice of statistical software, an instruction to use assistants for any work that they might normally use assistants for, and asking them to complete the analysis as they thought best, as though the research task had been their own idea, not trying

⁸This file included the state/year-level unemployment rate and labor force participation rate. Immigration policy flags were for policies for undocumented immigrants to get state drivers' licenses, to get college financial aid, to be banned from state public colleges, or to follow Omnibus immigration legislation that serves to increase the surveillance of immigration documentation. Additional indicators were for participation in E-Verify laws that require employers to verify immigration authorization, to limit E-Verify participation, participation in Secure Communities, and for participation in task-force or jail based 287(g) policies.

to match or not-match other researchers or guess what analyses the project organizers wanted to see. Once finished, they uploaded all of their code and data to a Sharepoint website, wrote a short description and interpretation of their results focusing on a single “headline” result, and filled out the research survey to report their results.

For Task 2, all of the previous instructions remained in place, but several were added to further specify the research design:

- There is a “treated” group that is comprised of all ethnically Mexican and Mexican-born individuals who are aged 26-30 on June 15, 2012 (recall that individuals must not have had their 31st birthday as of June 15, 2012 to be eligible for DACA).
- There is an “untreated” group that is comprised of people who would have been eligible for DACA, except that they were aged 31-35 on June 15, 2012.
- Researchers should estimate the effect of treatment by seeing how the 26-30 group changed from before treatment to after relative to how the 31-35 group changed (keeping in mind this is a repeated cross-section and not panel data).
- Researchers should attempt to estimate the effect for all individuals in the “treated” group and not, for example, estimate the effect only for men or only for women.
- The instructions specifically mention that researchers can, if they like, use covariates or account for differing trends to improve the comparability of the treated and untreated groups.

The task is otherwise unchanged for Task 2.

In Task 3, the instructions remain unchanged from Task 2, except that the data is provided directly instead of having researchers download data from IPUMS, omitting data from the year of 2012. In Task 3, project organizers cleaned the data, merged in the state policy data, created a variable indicating whether a given individual was in the “treated” or “untreated” group, limited the sample only to individuals in “treated” or “untreated,” and created simplified versions of variables like education. Researchers were instructed not to further limit the sample from this prepared data set, or to perform further extensive data cleaning.⁹

3.2 Recruitment and Attrition

In a many-analysts study, researchers who carry out the research task make up both the bulk of the author list and are the subject of inquiry, so their recruitment is a key feature of the study.

⁹There were three observations in the final cleaned data set that were missing values of the education variable. The final used sample in Task 3 sometimes differs by 3 across researchers, based on whether the analysis uses education and thus drops these individuals.

3.2.1 Researcher Qualifications

The goal of the project organizers was to make the set of researchers representative of the set of people who are producing the applied microeconomics literature. As such, recruitment criteria focused on identifying people who have produced applied microeconomic research, including potentially non-academic applied microeconomics research.

A given researcher was qualified for the project if they satisfied any one of the following criteria:

- They are academic faculty working in applied microeconomics.
- They are a graduate student **and** have a published or forthcoming paper in applied microeconomics.
- They hold a PhD **and** work in a job where they write non-academic reports using tools from applied microeconomics to estimate causal effects.¹⁰

Participation was not limited on the basis of country, career stage, or demographics such as sex, race, or sexual or gender identity.

3.2.2 Target Sample Size

An initial simulation-based power analysis assumed that each research task would have 5% less between-researcher variation in observed effects than the previous round and looked at the statistical power to detect a linear relationship between round number and the squared deviation of effects (variance of estimated effects across researchers). We found that we had 90% power to detect this effect if 90 researchers finished all tasks. We also found that, for comparisons of only two different research tasks, 90 researchers would give 85% power to detect a decline in variance from one stage to the next of 15% or more, a reasonable effect size given previous many-analyst studies.

We further assumed that attrition rates would be roughly 50%, which would suggest recruiting 180 eligible researchers to achieve adequate power. We revised that goal to 200 to account for our assumptions potentially being optimistic. Project organizers obtained funding to support payments to 200 researchers (see below).

¹⁰This qualification would allow, for example, employees of the World Bank, or people working in private sector research, to participate.

3.2.3 Recruitment and Incentives

Recruitment was advertised to potential researchers through three avenues: (1) social media posts on Twitter and LinkedIn, (2) emails to professional organizations including the Institute for Replication and the Committee on the Status of Women in the Economics Profession, and (3) emails to United States economics department chairs. For emails to departments heads, we gathered the list of all 286 economics departments listed in the U.S. News and World Report. We could locate email addresses for a front desk or (preferably) department chair for 264 of those departments. We emailed those 264 departments, asking for the message to be passed on to all faculty or just all microeconomics faculty.

The recruitment message described the project and its goals, and provided a link to a website that included further detail on project expectations and incentives for participation.¹¹ Researchers were told that if they completed all three tasks of the project, they would be offered a \$2,000 payment for up to 200 of the participants, and authorship on the eventual paper. The website included a link to a survey that asked questions related to eligibility for the project.

3.2.4 Participation and Attrition

Overall participation and attrition values are in Table 1. 362 people submitted applications for the project. 18.51% of these were found to be ineligible for the project. Most ineligible people were graduate students who did not yet have a forthcoming paper.¹²

Table 1: Participation and Attrition

Round	Participants	Attrition
Original Signup	362	18.51%
Assigned Task 1	295	47.80%
The first replication task	154	2.60%
The second replication task	150	2.67%
The third replication task	146	

This left 295 eligible participants. This is more than the 200 for which budget was available to pay the offered \$2,000 incentive. The 282 of these participants who had signed up by the original signup due date were put into a random order, and then the 13 late signups were put at the end of this order. Participants were given their place in the order, and informed that, among people completing all stages of the project, the first 200 in the order would be paid.

¹¹<https://nickch-k.github.io/ManyEconomists/>

¹²Data processing and analysis as well as table and figure creation for this paper were performed using the R packages `data.table`, `tidyverse`, `rio`, `fixest`, `car`, `modelsummary`, and `vtable` (Barrett et al. 2024; Wickham et al. 2019; Becker et al. 2023; Bergé 2018; Fox and Weisberg 2019; Arel-Bundock 2022; Huntington-Klein 2021).

Initial assumptions from the power analysis that attrition rates would be near 50% were almost exactly correct, with 49.49% of these initial 295 eligible researchers completing all three stages. Nearly all of the attrition occurred by the completion of Task 1. After 141 eligible researchers failed to complete Task 1, only a further 8 failed to complete Task 3. This means we have 146 researchers who completed all three research tasks, well above the goal of 90.

The high recruitment numbers and the fact that nearly all attrition occurs before Task 1 is complete allows us to evaluate the impact of the payment incentive. One potential concern with our incentive design is that payment and authorship are offered to anyone who completes all tasks, regardless of the quality of their work. We evaluate whether being guaranteed payment affects the probability of completing Task 1 using a regression discontinuity design. Someone randomly assigned to position 199 in the ordering is guaranteed payment if they complete all the tasks, while someone in position 201 may think they are likely to receive payment, but they are not guaranteed it.




Figure 1: Impact of Guaranteed Payment on Probability of Task 1 Completion

Figure 1 shows no meaningful effect of being guaranteed payment on the probability of completing Task 1. In additional results in the appendix, using a linear regression specification of the regression discontinuity design and the full range of the data (not including the late sign-ups) to maximize statistical power,¹³ we again find no statistically significant effect of being guaranteed treatment. This is suggestive that participants were not simply signing up in an attempt to get a \$2,000 payment for little effort.

¹³Use of the full range, rather than a bandwidth, is justified given that the running variable is randomly assigned aside from the late sign-ups. We also find no effect if we drop the late sign-ups from the regression discontinuity analysis.

3.2.5 Sample Characteristics

Tables 2 to 4 show the characteristics of the recruited sample, and how those characteristics changed with eligibility and attrition. Task 2 is omitted as an attrition stage since so few people dropped out between Task 1 and Task 2.

Table 2 shows that the majority of researchers were recruited via social media, with only about 9% coming from a department email, 4% from a professional organization email, and 9% from some other source (like word-of-mouth). Those recruited from another source were less likely to qualify for the study, and slightly less likely to finish, while those recruited from social media were most likely to qualify and finish. We also asked researchers how certain they were of their ability to finish the first task as well as the full set of tasks, on a scale of 1 to 100. Enrollees were about 90% confident in their ability to complete the full set of research tasks (although only about 50% did). Those who were more confident were slightly more likely to actually finish, and average confidence rates of those who did finish were about 92% instead of 90%.

Table 2: Researcher Recruitment Source and Completion Confidence

Round Variable	Original signup			Assigned task 1			Finished task 1			Finished task 3		
	N	Mean	SD	N	Mean	SD	N	Mean	SD	N	Mean	SD
Recruitment Source	347			285			150			142		
... Social media	270	78%		224	79%		124	83%		116	82%	
... Department email	31	9%		28	10%		13	9%		13	9%	
... Email of a professional organization	15	4%		10	4%		4	3%		4	3%	
... Other	31	9%		23	8%		9	6%		9	6%	
Certainty to Finish Task 1	355	90	11	292	90	10	153	92	8.4	145	92	8.3
Certainty to Finish Task 3	355	89	12	292	89	12	153	91	9.9	145	91	9.6

Table 2 shows the professional experience of enrollees. While graduate students were considered eligible for the project as long as they had a published or forthcoming paper, the majority of eligible researchers (83%) had PhDs. PhD holders were also more likely than other eligible researchers to complete all three tasks.

These PhDs are split across faculty (62%) and other non-faculty researchers (22%), both of which were more likely than graduate students to finish all three rounds. Note that the researchers in these categories who do not hold PhDs were either people who had been hired to faculty roles without holding PhDs (such as ABDs, or people in a faculty position requiring only a Master’s degree), or people with Master’s degrees in non-faculty research positions who had published academic papers (some of whom were still graduate students).

Most of the researchers had at least one published paper, and researchers with 6+ papers were more likely than others to complete all three research tasks. Those with “No Academic Papers” are non-academic researchers who produce work not intended for academic journal publication. Those with “No Published Academic Papers” have papers that are forthcoming, or are faculty who only have working papers and no publications.

The set of researchers in the study generally do not work in the specific subfield that the research task is in. The research task is similar to many studies done across all of applied microeconomics, but specifically is on the topics of labor and immigration. About a third of the enrollees had done research in either immigration or labor previously, and these researchers were somewhat more likely to complete all three tasks. No researchers enrolled who had previously worked in both immigration and labor.

Table 3: Researcher Professional Experience

Round	Original signup		Assigned task 1		Finished task 1		Finished task 3	
Variable	N	Percent	N	Percent	N	Percent	N	Percent
Degree	360		295		154		146	
... No graduate school	3	1%	0	0%	0	0%	0	0%
... Some Grad School	14	4%	5	2%	3	2%	2	1%
... Master's degree	78	22%	44	15%	17	11%	17	12%
... Prof. Degree	3	1%	1	0%	0	0%	0	0%
... PhD	262	73%	245	83%	134	87%	127	87%
Occupation	361		295		154		146	
... Faculty	191	53%	182	62%	99	64%	98	67%
... Grad. Student	69	19%	36	12%	13	8%	12	8%
... Other	14	4%	11	4%	5	3%	3	2%
... Other Researcher	87	24%	66	22%	37	24%	33	23%
Research Experience	361		295		154		146	
... 1-5 Papers in Applied Micro	162	45%	152	52%	74	48%	70	48%
... 6+ Papers	104	29%	102	35%	58	38%	57	39%
... No Academic Papers	17	5%	4	1%	3	2%	3	2%
... No Published Academic Papers	78	22%	37	13%	19	12%	16	11%
Field	333		270		145		138	
... Immigration & Labor	0	0%	0	0%	0	0%	0	0%
... Immigration	8	2%	6	2%	4	3%	4	3%
... Labor	102	31%	85	31%	49	34%	47	34%
... Neither	223	67%	179	66%	92	63%	87	63%

While the research tools used are not listed in Table 3, we did check the programming languages used by researchers. The most common language was Stata, with 109 researchers performing their work solely in Stata. 33 used R, and one researcher used both R and Stata. Less common were Python and SPSS, with one researcher each.

Table 4 shows the demographics of the researcher sample. The eligible sample was just under 80% male and more than 55% white, and both percentages grew by the conclusion of Task 3, with the white share growing significantly to 66%. The 80% male figure is similar to the share male found for faculty at a selected set of top economics departments in 2017 by Lundberg and Stearns (2019), and among all actively publishing economists in 2019 by Card et al. (2022). A small share reported being LGBTQ+, and this share remained constant over all rounds of the research tasks. An additional form of demographic difference is geographic. About half

of the sample was situated in the United States, and about half was from another country.¹⁴ The representativeness of the racial mixture is difficult to assess for this reason; 66% white would be low if the entire sample were from the United States (Stansbury and Schultz 2023), but it is unclear what the population rate is in a 50% US/50% other location sample.

Table 4: Researcher Demographics

Round	Original signup		Assigned task 1		Finished task 1		Finished task 3	
Variable	N	Percent	N	Percent	N	Percent	N	Percent
Gender	359		294		154		146	
... Female	81	23%	64	22%	28	18%	26	18%
... Male	274	76%	230	78%	126	82%	120	82%
... Non-binary / third gender	1	0%	0	0%	0	0%	0	0%
... Prefer not to say	3	1%	0	0%	0	0%	0	0%
Race	360		294		154		146	
... White	188	52%	164	56%	100	65%	97	66%
... Asian	79	22%	60	20%	25	16%	25	17%
... Black or African American	27	8%	21	7%	4	3%	4	3%
... Hispanic	25	7%	19	6%	10	6%	9	6%
... Other or Multiracial	41	11%	30	10%	15	10%	11	8%
LGBTQ+	360		294		154		146	
... Yes	18	5%	14	5%	7	5%	7	5%
... No	323	90%	268	91%	137	89%	129	88%
... Prefer not to say	19	5%	12	4%	10	6%	10	7%

One researcher did complete all three research tasks, and appears in the above tables, but their work has been removed from the results that follow in the rest of the paper, as due to a misunderstanding of the instructions, their work did not attempt to estimate the effect of DACA on the probability of employment.

As a whole, the sample largely reflects the group of people who publish work in applied microeconomics. The sample is skewed towards the United States, which is partially driven by the emails sent to US economics departments, the fact that the project was advertised and carried out in English, and the fact that the project organizers are in the United States and advertised the project using their own social media. Given that caveat, the makeup of the sample appears to be fairly similar to the makeup of the profession itself, although this is difficult to verify for some demographics.

¹⁴Exact figures are not given for geography, and crosstabulations across geography are not given, because non-geographic demographic information comes from a survey where we acquired permission to share aggregate figures. Geographic information, on the other hand, comes from researcher payments information, for which we did not request permission to share responses.

4 Results

This section demonstrates the variation in effects and methods across researchers and conditions, both demonstrating that variation and attempting to explain it.

Importantly, these results are derived from the survey responses that researchers gave about their findings and the choices made. Project organizers did not cross-reference survey responses against researcher code to ensure that their code was accurately reflected in the survey, except in a small number of cases where the survey response could not be interpreted. This means that the variation presented here represents the variation in how researchers would plan to implement the research task if they were doing it independently, and what a reader would see as the description of a study in a published version of their work. Any variation between researchers that occurs as a result of coding error or a research report that misrepresents what a researcher actually did will not be reflected here, but could be the subject of a future investigation.

4.1 Variation in Effects and Sample Sizes

Figure 2 and Table 5 show the distribution of estimated effects across all researchers. The effect distributions are shown in two ways: unweighted and using inverse-standard-error weights.¹⁵ Several data points are dropped from the weighted analysis for researchers who did not report standard errors or reported 0. Other missing values are researchers who did not respond to a given question.

In Task 1, the mean estimated effect of DACA eligibility on the probability of working full-time was .053 unweighted or .044 weighted. In both cases these means are pulled upwards by high top-end estimates and are above the 75th percentiles. Median estimates were .030 unweighted or .026 weighted. Even in Task 1, with a large amount of freedom afforded, researchers found a reasonable amount of agreement in the effect sizes outside of the tails, with the 25th to 75th percentile ranges of the effect being .014 to .051 unweighted, an inter-quartile range (IQR) of .037, or 3.7 percentage points in the effect, or .012 to .043 weighted, an IQR of .031. The use of weights narrows the distribution of effects: researchers reporting smaller standard errors also reported estimates that were more similar to each other, which was also the case in Tasks 2 and 3.

Our preregistration plan details that we planned to give descriptive characteristics of the results, as we do here, and also implement a Levene test on whether the variance between

¹⁵The use of inverse-standard-error weights is not preregistered but follows meta-analytic standards, reducing the influence of estimates that may be outliers due to being estimated with a highly-noisy method, under the suggestion of Auspurg and Brüderl (2023). Weights are truncated at the 95th percentile (200, or a standard error of .005) so as to avoid any single researcher having too much influence on results. Skipping the truncation shows more agreement because a few researchers with very small standard errors make up a significant share of the weighted sample.

researchers declined. We do not reject at the 95% level the null of no change in variance from any stage to any later stage (including a comparison of each task to its revision stage, and comparing each main task to later main tasks), with the lowest p-value of 0.197 coming from the comparison of Task 1 to Task 1 Revision.

Task 2 is somewhat odd in that it shows less agreement than Task 1 despite giving researchers less freedom. The IQRs increase to .043 unweighted or .040 weighted. Further, the effect distributions are somewhat bimodal, especially when weighted. One of these modes appears to be researchers reporting effect estimates of a similar level to those in Task 1, and others reporting effect estimates similar to what would later be found in Task 3.

Moving all the way to Task 3, agreement considerably increases between researchers. The 25th and 75th percentile effects are .031 and .058 unweighted (IQR .027), and .036 and .060 weighted (IQR .024). The bimodality from Task 2 is still there, but with much more agreement and density at the higher mode. From Round 1 to Round 3 we see considerable increases in agreement between researchers.

Taking only the effect distributions as a baseline, we see that, at least in this application, researchers in general report fairly similar, although certainly not identical, effect estimates on average, but there are some extreme outlier estimates as well. We may also take this to mean that providing pre-cleaned data, as in Task 3, led to a strong increase in researcher agreement. Specifying further the research question and design, however, as in Task 2, led to somewhat less agreement. The odd result for Task 2 and its proper interpretation will be investigated further in Section 4.6.

Table 5 also shows the reported standard errors. Reported standard errors increase significantly from round to round, driven largely by the specification of the research design, which for many researchers considerably narrowed the sample they were supposed to use. This can also be seen in Figure 3, where the distribution of effects narrows a little across rounds, but confidence intervals increase considerably. Throughout, while there is general agreement on effect size in the middle of the distribution, researchers vary in whether the reported effect is statistically significant, with 78%, 60%, and 64% reporting results that were statistically significantly different from 0 in Tasks 1, 2, and 3, respectively.

We can compare average standard errors against the variation in effects between researchers to get a sense of how much effect variability is omitted by only considering a reported standard error, as in Huntington-Klein et al. (2021) or Menkveld et al. (2021). Comparing the standard deviation of weighted effects against the average standard error gives ratios of 4.84, 2.23, and 1.75 for Tasks 1, 2, and 3, respectively. Huntington-Klein et al. (2021) used a single round and allowed full researcher freedom, and found a range of 3-4 for this ratio, below what we find for the full-freedom Task 1. If we instead compare weighted IQR to average standard errors we get ratios of 1.63, 1.29, and .41. This is partially driven by increasing agreement over rounds, but is also driven by the addition of restrictions that reduce sample size and thus increase average standard errors. These figures suggest that a reported standard error considerably understates the estimate uncertainty that we should acknowledge when including researcher




Figure 2: Distributions of Reported Effect Sizes

variation. However, these figures also demonstrate a flaw with these ratios, introduced in Huntington-Klein et al. (2021), as a metric for researcher-induced uncertainty: in that they are sensitive to the estimate precision one might be expected to get given the research task. Researcher variation does not scale with that uncertainty.


Table 5: Distribution of Reported Effects and Sample Sizes

Variable	N	Mean	SD	Min	Pctl. 25	Median	Pctl. 75	Max
Round: Task 1								
Effect Size (Unweighted)	145	0.053	0.095	-0.049	0.014	0.030	0.051	0.660
Effect Size (Weighted)	138	0.044	0.092	-0.049	0.012	0.026	0.043	0.660
Standard Error	139	0.019	0.055	0.000	0.005	0.007	0.013	0.460
Sample Size	145	828,318	3,056,037	681	61,600	179,960	356,787	29,536,580
Treated-Group Size	141	96,395	648,493	270	17,950	34,435	52,581	7,727,201
Round: Task 2								
Effect Size (Unweighted)	145	0.044	0.100	-0.390	0.015	0.032	0.058	0.850
Effect Size (Weighted)	141	0.046	0.069	-0.090	0.018	0.034	0.058	0.850
Standard Error	141	0.031	0.078	0.001	0.010	0.014	0.020	0.744
Sample Size	144	157,006	1,065,593	6,196	18,981	25,414	48,125	12,609,847
Treated-Group Size	140	31,948	221,175	3,519	5,953	11,157	15,832	2,627,183
Round: Task 3								
Effect Size (Unweighted)	145	0.045	0.101	-0.810	0.031	0.050	0.058	0.650
Effect Size (Weighted)	142	0.062	0.103	-0.810	0.036	0.051	0.060	0.650
Standard Error	144	0.059	0.268	0.000	0.015	0.018	0.026	2.747
Sample Size	145	16,904	1,756	7,833	17,379	17,382	17,382	17,832
Treated-Group Size	129	9,433	3,008	11	5,149	11,382	11,382	17,383

Table 5 also shows summary statistics for reported standard errors, both overall and for the treated group. These distributions are also shown in Figure 4 and Figure 5.

In Figure 4 we see a huge amount of variation in the reported sample size used in Task 1, noting that the x-axis is on a log scale. The 25th and 75th percentiles of reported sample sizes ranging from 61,600 to 356,787, with some researchers using millions of observations.¹⁶ For Task 2, which specified in the instructions the treated and comparison groups to use, variation reduces considerably, although the 75th percentile (48,125) is still double the 25th (18,981), and there are still some researchers using millions of observations. Task 3 is not shown in the graph because the sample is pre-specified, with the only meaningful variation being whether or not the researcher dropped three rows of data with missing education values, and a few outliers reporting lower numbers. The lower sample sizes for this question are due to researchers who skipped it because they assumed the answer was obvious.

¹⁶Keep in mind that for Task 1, there was not a specified control group, so a researcher may decide to use the entire ACS sample in the analysis, including people very unlike the eligible group in the sample. In Task 2, the instructions specified a treated and comparison group, but some researchers may have different samples than in Task 3 either due to error, or because they included people other than the treated and comparison group in their sample to improve precision, and used their model to compare those groups more directly.



95% CI reconstructed from effect size and SE, even if asymmetric CI was reported. Visible range limited to $(-.05, .15)$.

Figure 3: Specification Curve for All Reported Estimates




Figure 4: Distributions of Reported Sample Sizes

Variation in the reported size of the treated group in Figure 5 is affected somewhat by researcher confusion in responding to the survey question. The survey question about treated-group size instructed researchers not to count individuals eligible for DACA as treated for the purposes of this question if they were in a pre-DACA year. However, many researchers counted these individuals as treated anyway, leading to variation in the Task 3 distribution, even though every researcher is at this point working with the same eligibility indicator.




Figure 5: Distributions of Reported Treated-Group Sizes

Aside from this issue, we see that the imposition of a shared definition for the treated group reduced the IQR for the size of the group from 34,631 in Task 1 to 9,879 in Task 2. Theoretically, however, since there was a shared definition of the treated group in Task 2, there should be no more variation in this variable in Task 2 than in Task 3. This indicates that not all instructions were implemented in the same way across researchers, which will be explored further in Section 4.4. Despite a shared understanding of who was eligible for DACA and who should be in the treated group, only a shared data preparation that implemented these rules for people led to sharp agreement in the size of the treated-groups sample.

4.2 Peer Review

This section evaluates the impact of peer review on the later work performed by a researcher. The structure of peer review in this study is that, following each main task, 2/3 of the researchers are randomized into pairs that produce a peer review report of the other's work,

while the remaining 1/3 do not receive or perform peer review. Then, researchers have an opportunity to revise their work.

Revision is optional, and relatively few researchers (about 30 per task) chose to revise their work after receiving peer review. As such, we mostly look at the impact of peer review on the work performed in subsequent main tasks. However, we can check if revision tended to lead towards more agreement overall. In Table 6, we show the variance of the entire sample of reported effects post-revision, replacing each researcher’s reported task effect with its revision, if they revised their work, and compare variance among the peer-reviewed group to the non-peer-reviewed group. At no point is this difference statistically significant, and it is inconsistent which group had the lower post-revision variance. Taking only the subgroup of actually-revised estimates, we see that these tended to agree with each other more than the group as a whole did in Tasks 1 and 2, but this is also inconsistent, with greater variance than the whole group in Task 3 and when pooling over all tasks (after subtracting by-task means).

Table 6: Post-Revision Variance by Peer Review

Task	Unreviewed Variance	Reviewed Variance	Levene Test p-value	Revised Variance
Task 1	0.002	0.012	0.173	0.001
Task 2	0.009	0.004	0.571	0.002
Task 3	0.001	0.008	0.210	0.015
Pooled	0.004	0.008	0.219	0.005

In general, the mechanisms by which peer review might be expected to change a researcher’s work in normal journal submissions include both that researchers might find peer review comments helpful and incorporate them into their work, and that researchers are required by the journal submission process to incorporate most reviewer comments. In this study, our peer review process can only capture the first of these mechanisms.

Figure 6 shows the distribution of effect sizes estimated by those who did, and did not, engage in peer review in each round. The left column of graphs shows the effects reported in each task before researchers were assigned to peer review, and the right column shows the effects reported in the follow-up task, comparing those either assigned or not assigned to peer review in the previous task. As we might expect given random assignment, effect distributions are fairly similar pre-review between the review and non-review groups, with a slightly narrower distribution of effects for researchers about to be reviewed.

We also see that these groups have very similar effect distributions in their follow-up task. Neither group has a considerably narrower distribution than the other. Levene tests for differences in follow-up round effect variance between the peer-reviewed and non-peer-reviewed groups show p-values of 0.846 and 0.788 in Tasks 2 and 3, respectively, or 0.999 when pooling

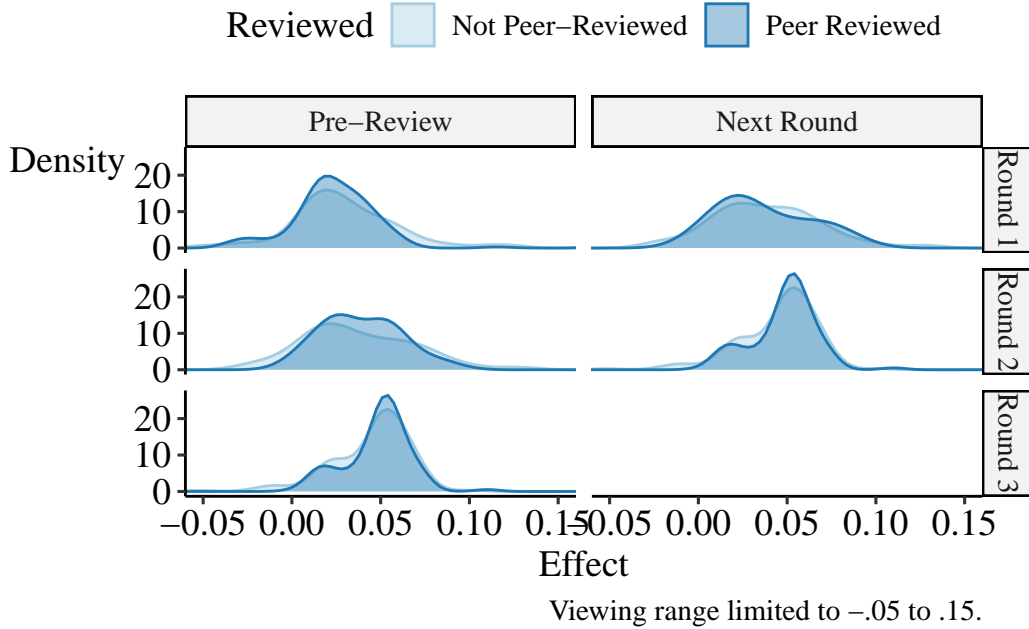


Figure 6: Distributions of Reported Effect Sizes

the two tasks. This is not strong evidence in favor of the idea that peer review might drive agreement between researchers due to the receipt of feedback.

Figure 7 explores the possibility that peer review might not make the peer-reviewed group as a whole more similar, but rather just make someone more similar to their specific reviewer. We calculate the absolute difference in effects between each reviewer pair, in the task they perform before reviewing (left column), in the follow-up task (middle column) and comparing your follow-up task against your reviewer’s result this round (right column), with the right column representing the possibility that a researcher may select an analysis so as to produce a result more similar to the one they saw in the previous round. The distributions of absolute differences for non-reviewed researchers are generated as a null distribution by matching every non-reviewed researcher to every other non-reviewed researcher and calculating all absolute differences.¹⁷

In Figure 7 we see some of the anticipated effect of peer review for Task 1. Before review (left column), absolute differences between review pairs were more likely to be large than differences between non-review pairs. But by the follow-up in Task 2, both groups were similar, potentially suggesting that peer review reduced the large absolute differences. In the right column, the

¹⁷This null distribution represents the distribution of absolute differences among people who did not actually experience peer review. Notably, each non-reviewer is matched multiple times in this approach, instead of just once for reviewers. However, matching the non-reviewers only once to a single random pair just produces a noisier version of this all-matches null distribution. Averaging the single-random-match approach over many random single matches produces the same null distribution.

unreviewed group still shows lower differences, but to a smaller degree. However, none of this holds up in Task 2. Again the actual review pairs started out with larger differences, and those differences shrank for both groups by Task 3, but by the same degree. Appendix Table 20 shows that average absolute differences grew by a statistically significant .051 more for the unreviewed group than for the reviewed group in Task 1, but that this effect reverses to a statistically significant .029 in favor of the unreviewed group for Task 2. This is not consistent strong evidence of peer review making a researcher more like their reviewer as the result of feedback.

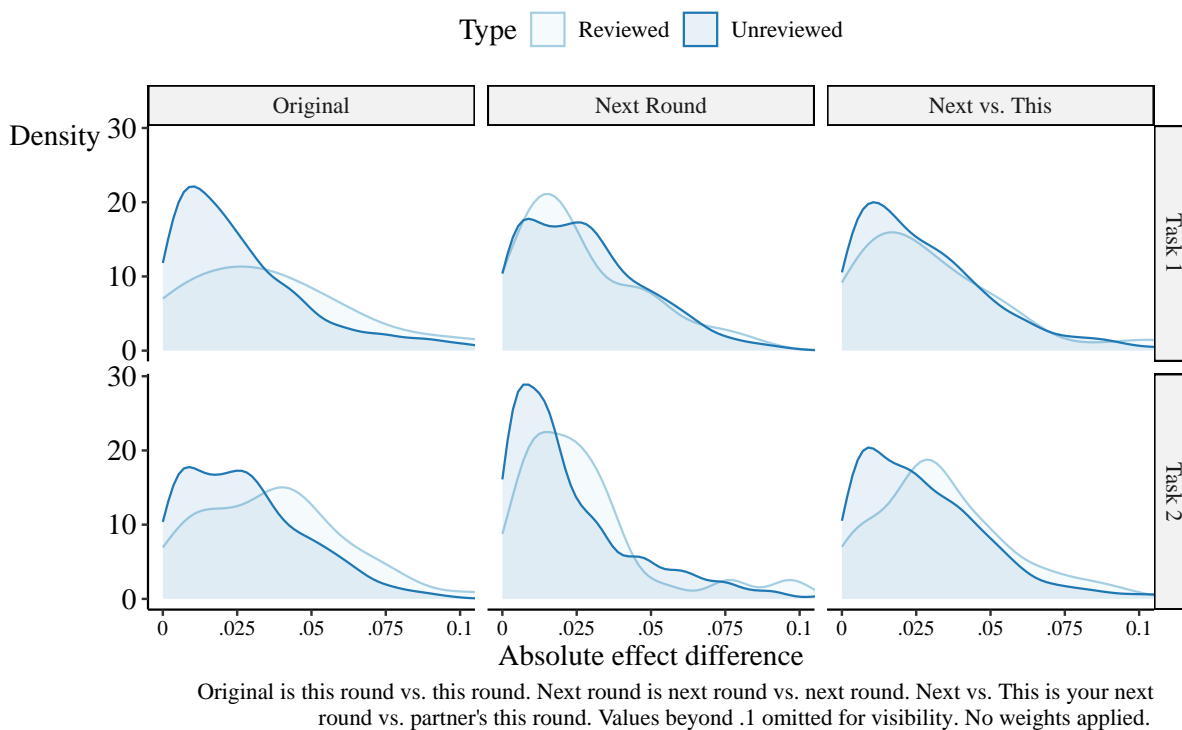


Figure 7: Comparisons of Effect Sizes vs. One's Reviewer

4.3 Analytic Choices

The next two sections examine the different choices that researchers made, both to demonstrate the variation in ways that different researchers chose to carry out the research tasks, and to relate those choices to differences in outcomes.

Table 7 shows the different choices made in estimating the effect of DACA on the probability of employment across all three tasks, in particular the estimator chosen, the use of ACS sampling weights provided by IPUMS, and the choice of standard error adjustment.¹⁸ The

¹⁸For most researchers, these choices did not change over the tasks, and so we just present the overall view.

dependent variable of interest, working full-time or not, is binary. However, as is generally standard in applied microeconomics, linear regression was the most common estimator used, with 82% of entries. 13% used logit or probit regression instead. Notably, many researchers used linear regression as a means of implementing a fully saturated (or nearly fully saturated) difference-in-differences design, in which case the downsides of linear probability models are muted. Other researchers mostly used a matching estimator (sometimes combined with linear regression) or one of several newly-introduced estimators for difference-in-differences designs, like Callaway and Sant’Anna (2021).

The use of sample weights was relatively uncommon, despite their use being advised with survey data like ACS. Only 25% of task completions mentioned the use of weights or any of the standard ACS weight variables.

There was considerable variation across researchers in the selection of standard error adjustment. A slim majority of researchers applied clustered standard errors in some way, but clustered at different levels: state, state/year, or according to a survey clustering indicator like Strata or some other variable. A further 17% of submissions used heteroskedasticity-robust but not cluster-robust standard errors.

Table 7: Estimation Methods

Variable	N	Percent	Variable	N	Percent
Method	437		S.E. Adjustment	438	
... Linear Regression	358	82%	... Cluster (State)	118	27%
... Logit/Probit	57	13%	... Cluster (State & Year)	58	13%
... Matching	11	3%	... Cluster (ID/Strata/Other)	65	15%
... New DID Estimator	7	2%	... Het-Robust	76	17%
... Other	4	1%	... Other/Bootstrap	23	5%
Weights	438		... None	98	22%
... No Sample Weights	329	75%			
... Sample Weights	109	25%			

This table shows details on estimation, not research design. "Difference-in-differences" implemented with linear regression, for example, counts here as linear regression.

Table 8 shows the average rate of inclusion of covariates across all three tasks. These can be read as the share of researchers who included the covariate, with the exception of "Other", which allows each researcher to have multiple "Other" controls. The most common covariates included are shown, with the exception of indicators for "eligible for DACA" or "in a post-DACA period", as these are considered part of the core research design rather than covariates. Variables are included here regardless of the functional form used to include them.

The most common included controls were for state, year, age, and sex, which were included as covariates for more than 50% of researchers in all three tasks. There was heavy variation

in the sets of included covariates, however. In Task 1, for example, there are ten covariates with rates between .2 and .8, meaning that at least 20% of the researchers made a different decision on inclusion of the covariate than the majority. There are four covariates in the 40-60% range, meaning that the researchers were almost evenly split on whether or not to include the covariate. These rates did not change much by Task 3.

Across all rounds, in which there were 435 submitted research tasks, there were 333 different unique sets of included covariates after “Other” covariates are excluded. 64% of submissions had a set of covariates that was unique for the task. 17% shared a covariate set with one other person in that task, 12% shared with two or three other people, and only the those with no controls shared with more than three other people.

Table 8: Average Rate of Covariate Inclusion Across Rounds

Control	Task 1	Task 2	Task 3
Age	0.62	0.57	0.52
Age at Migration	0.18	0.14	0.14
Age in 2012	0.05	0.04	0.08
Continuous Years in USA	0.13	0.13	0.12
Education	0.48	0.50	0.51
English Speaker	0.17	0.17	0.23
Labor Force Participation Rate	0.22	0.17	0.20
Marital Status	0.10	0.13	0.12
Race	0.24	0.22	0.28
Sex	0.63	0.64	0.72
State	0.62	0.64	0.64
State Policy Variables	0.25	0.21	0.23
Unemployment Rate	0.32	0.27	0.30
Year	0.68	0.60	0.57
Year of Migration	0.14	0.13	0.11
Other	0.66	0.63	0.63
None	0.07	0.07	0.06

There was very little agreement across researchers on the exact set of appropriate controls, or the inclusion or exclusion of any given control (aside from those very rarely included). Did these choices impact the effect estimates? Not by much. Table 9 shows the average effect among researchers who included a given control, pooling all three tasks. The mean reported effects differ by only .023 percentage points comparing the covariate included in analyses with the highest average effect estimates (Continuous Years in the USA) against the lowest (Race). This likely overstates the impact of covariate selection here, as selecting the highest vs. lowest after estimates are known will bias us towards a larger difference from noise alone.

There do not appear to be major differences in the average reported standard errors either, or in the standard deviation of the effect distribution among reserachers including that covariate.

Table 9: Estimated Effects by Control-Variable Inclusion

Control	N	Effect	Mean SE	Effect SD
Continuous Years in USA	55	0.054	0.035	0.123
Age	248	0.048	0.025	0.094
Year of Migration	55	0.048	0.033	0.112
Marital Status	51	0.047	0.016	0.071
Sex	289	0.046	0.027	0.101
Age at Migration	67	0.045	0.022	0.067
None	28	0.045	0.060	0.133
State	275	0.045	0.025	0.089
Year	267	0.045	0.026	0.094
Education	216	0.042	0.017	0.061
Other	278	0.040	0.038	0.086
Age in 2012	24	0.037	0.026	0.042
State Policy Variables	99	0.037	0.033	0.108
Unemployment Rate	128	0.036	0.033	0.097
Labor Force Participation Rate	86	0.035	0.041	0.115
English Speaker	82	0.034	0.046	0.100
Race	106	0.031	0.032	0.092

While the inclusion of a given common control variable does not strongly predict an estimated effect, in Table 10 we look at the most common covariates and examine whether their functional form meaningfully affects the estimated effect. The selection of functional form explained more variation in average estimated effects than the inclusion of covariates did, at least in this context. For both age and the State/Year controls, the difference between the highest-average-effect functional form variants and the lowest, in both cases comparing a linear control against a fixed effect, was greater than the difference between highest and lowest among covariates included.

Table 10: Estimated Effects by Functional Form of Control Variable

Category	Control	N	Effect	Mean SE	Effect SD
AGE	Linear Age	164	0.058	0.024	0.107
AGE	Age FE	36	0.024	0.040	0.022
AGE	Age Quadratic	33	0.035	0.015	0.089
EDUC	Linear Education	122	0.040	0.016	0.066
EDUC	Education FE	32	0.047	0.021	0.033

Category	Control	N	Effect	Mean SE	Effect SD
EDUC	Education Transform	61	0.045	0.017	0.064
STATE/YEAR	Linear Year	79	0.044	0.037	0.140
STATE/YEAR	Year FE	103	0.047	0.026	0.062
STATE/YEAR	State FE	155	0.046	0.031	0.102
STATE/YEAR	State FE x Year FE	56	0.037	0.018	0.027
STATE/YEAR	State FE x Linear Year	23	0.061	0.017	0.133

These specific findings about the impact of choices on effects - that the inclusion of different covariates did not have a major impact on estimated effects, or that the choice of functional form had a greater impact than the selection of covariates - should not be expected to generalize, and is specific to this research task. However, what we can take from this section is that there is heavy variation across researchers in what they believe the appropriate set of covariates to be and, for a given covariate, what the appropriate functional form is. We can also see that, in the case of this particular study, these decisions, while varied, did not fully explain the variation in effects between researchers.

4.4 Sample Limitations

In this section we examine the ways in which researchers defined their analytic samples, as well as defined the treated group, and any comparison group included in the data analysis that was not treated. These data are derived from researcher responses to a survey about their work, in which they were asked to describe how they limited the size of their sample before downloading it from IPUMS and which variables they used to further drop observations from the sample before analysis. For example someone might say that they only kept observations for which `HISPAN == 1` (the observation is Hispanic-Mexican), along with other restrictions. These two responses are combined to define the full analytic sample.

Researchers were also asked to describe how they defined the treated group who was affected by DACA. For example someone might say that only those with `CITIZEN == 3` (non-citizens) were eligible, along with other restrictions. These were combined with the analytic-sample restrictions to make the full treated group definition. Finally, researchers were similarly asked to describe the conditions that defined someone who would be included in analysis but not be treated by DACA, which were combined with the analytic-sample restrictions to make the full comparison group definition.

Researchers were asked to specify these conditions as precisely as possible to match what they did, and to use IPUMS variable names in their descriptions. The project organizers coded these into a set of boolean conditions defining the overall sample, treated group, and comparison group for each researcher, in some cases reviewing the code directly or asking researchers to clear up uncertainty where survey responses were unclear, but in general taking the researcher

survey responses at their word. For these analyses, Task 3 is omitted because the analytic sample is defined for all researchers.

Table 11 looks purely at the number of distinct variables referenced in the sample limitations, regardless of what they are. In Task 1, the typical researcher used five variables to define their analytic sample, and an additional four to define their treated group. In Task 2, where inclusion criteria were shared, both of these numbers increased, but there was still considerable variation, with the 25th and 75th percentiles using 3 and 10 variables to define their full sample. Definition of the treated group was more shared, with the 25th and 75th percentiles using 9 and 12 variables, respectively.

Table 11: Number of Variables Referred to in Sample Limitations

Variable	N	Mean	Std. Dev.	Min	Pctl. 25	Pctl. 50	Pctl. 75	Max
Round: Task 1								
Whole Sample	145	5.0	2.7	0.0	3.0	5.0	7.0	14.0
Treated Group	145	8.6	2.1	3.0	7.0	9.0	10.0	17.0
Untreated Group	145	8.0	2.7	1.0	7.0	8.0	10.0	17.0
Round: Task 2								
Whole Sample	145	6.8	3.9	0.0	3.0	8.0	10.0	15.0
Treated Group	145	10.3	2.3	3.0	9.0	10.0	12.0	16.0
Untreated Group	145	10.2	2.3	3.0	9.0	10.0	12.0	16.0

As for how those variables are actually used, Table 12 shows how these variables were implemented as sample restrictions. Importantly, these reflect the survey responses given by researchers. Some of the “none” responses in Task 2 in particular reflect researchers who did use the variable in some way but did not report it in their description of results.¹⁹

We see a huge amount of variety in the ways these variables were used, including in Task 2 for the treated-group definition, where there is a correct answer according to the instructions (and similarly a correct answer for some variables in the Task 1 treated-group definition).²⁰ For

¹⁹Other notes of interest for reading the table: (a) “Multistep condition” refers to cases where the variable is included, but only as a part of a complex boolean statement involving many variables. These most commonly appeared in definitions for the comparison group, which are not in the table, in the format of “fails any one of the following set of DACA eligibility requirements.” and (b) for Education/Veteran status, recall that the mention of this eligibility requirement was omitted from the Task 1 instructions, which explains why very few researchers used these variables to define their samples or treated groups in Task 1.

²⁰Keep in mind also that these are recoded versions of the actual survey submissions sent in by researchers. The survey question asked respondents to use IPUMS variable names to describe their choices. So for example the individuals reporting that they used only high school graduates or *non-veterans*, instead of veterans as per the instructions, likely did not intentionally choose to use non-veterans and write “I chose to use non-veterans in my sample” in their response but rather they wrote “VETSTAT == 1,” which indicates “non-veteran”, perhaps based on a misunderstanding of the IPUMS documentation (veterans are VETSTAT

each variable, the most-common option, listed at the top, is the “correct” answer for defining the treated group, with two exceptions: (a) for Citizenship, there is a second justifiable answer in “Non-Citizen or Naturalized After 2012.” These immigrants would have been eligible for DACA in 2012, but would not be eligible for DACA as of the time they were surveyed, so they would have received a partial “dose” of DACA, which could justifiably be included or excluded, and (b) for Years Continuous in USA, where DACA guidelines require that the immigrant have lived *continuously* in the United States for five years as of 2012. Most researchers used only year of immigration being before 2007 to satisfy this criterion, but others used the YRSUSA set of variables which specifically track living continuously in the country.

For all other variables besides Years Continuous in USA, the option matching the instructions was the most common, but we also see plenty of variation. We also see considerable variation for the columns in which there is not a clear “correct” option, like the analytic sample definition. No single way of applying any variable was used by more than 84% of the sample in any case. One interesting feature is the use of both “< 16” and “≤ 16” for age at migration, and “< 2007” and “≤ 2007” for year of migration. For year of migration, the two are similarly popular. Also interesting is the distinction between researchers using age defined in years to determine eligibility vs. age defined in quarters, which makes a difference given that eligibility is based on age specifically in June 2012.

Showing the impact of these choices on estimated effects is difficult since, aside from the most-common option, any specific alternative does not have enough people using it to make a reasonable comparison. However, we show estimated effects and, for analytic-sample restrictions, analytic sample size by sample limitation choice in Appendix Tables 21 for Task 1 and Table 22 for Task 2. There are large differences in estimated effects and sample sizes across many of these different sample restriction choices, but in many cases these comparisons are based on very small samples.

The two comparisons for which an alternative was common enough to compare are for the YRSUSA inclusion and the use of “< 2007” vs. “≤ 2007” for year of migration, which are shown in 13 for Task 1 and Table 14 for Task 2. For both of these, the relationship between these choices on effects varies from negligible to a several percentage-point difference associated with a single sample restriction change, a fairly minor one in particular for “< 2007” vs. “≤ 2007”. Effect differences are larger in Task 2. However, in Task 1, even though estimated effects are similar, sample sizes are considerably larger for the less-restrictive option, and so reported standard errors would be lower, and statistical significance more likely.

4.5 Researcher Characteristics and Effects

In this section we evaluate the relationship between researcher characteristics and the effects they reported. As in our preregistration, analysis in this section is performed in a multiple-

== 2).

Table 12: Sample Restriction Methods

Round/Sample Variable	Task 1 All		Task 1 Treated		Task 2 All		Task 2 Treated	
	N	Percent	N	Percent	N	Percent	N	Percent
Hispanic	145		145		145		145	
... Hispanic-Mexican	93	64%	106	73%	101	70%	112	77%
... Hispanic-Any	7	5%	8	6%	7	5%	8	6%
... Hispanic-Mex or Mex-Born	5	3%	6	4%	0	0%	1	1%
... Multistep Condition	2	1%	2	1%	1	1%	1	1%
... None	38	26%	23	16%	36	25%	23	16%
Birthplace	145		145		145		145	
... Mexican-Born	91	63%	88	61%	96	66%	95	66%
... Hispanic-Mex or Mex-Born	4	3%	2	1%	0	0%	1	1%
... Non-US Born	2	1%	1	1%	3	2%	0	0%
... Central America-Born	1	1%	1	1%	1	1%	1	1%
... None	47	32%	53	37%	45	31%	48	33%
Citizenship	145		145		145		145	
... Non-Citizen	71	49%	109	75%	88	61%	118	81%
... Foreign-Born	3	2%	3	2%	1	1%	0	0%
... Non-Cit or Natlzd post-2012	1	1%	4	3%	2	1%	5	3%
... Citizen (various)	1	1%	3	2%	1	1%	1	1%
... Multistep Condition	2	1%	1	1%	0	0%	0	0%
... Other	7	5%	9	6%	3	2%	7	5%
... None	60	41%	16	11%	50	34%	14	10%
Age at Migration	145		144		145		145	
... < 16	12	8%	85	59%	52	36%	99	68%
... <= 16	5	3%	17	12%	7	5%	14	10%
... Other	30	21%	24	17%	12	8%	14	10%
... None	98	68%	18	12%	74	51%	18	12%
... > 16	0	0%	0	0%	0	0%	0	0%
... Any Age	0	0%	0	0%	0	0%	0	0%
... Multistep Condition	0	0%	0	0%	0	0%	0	0%
Age in June 2012	145		145		145		145	
... Year-Quarter Age	25	17%	122	84%	78	54%	129	89%
... Year-Only Age	13	9%	9	6%	9	6%	11	8%
... None	107	74%	14	10%	58	40%	5	3%
Year of Immigration	145		145		145		145	
... < 2007	12	8%	39	27%	30	21%	45	31%
... <= 2007	10	7%	50	34%	24	17%	44	30%
... < 2012	2	1%	1	1%	0	0%	2	1%
... <= 2012	2	1%	2	1%	1	1%	3	2%
... >= 2007	1	1%	1	1%	0	0%	1	1%
... Any Year	7	5%	3	2%	2	1%	2	1%
... Multistep Condition	2	1%	1	1%	0	0%	2	1%
... Other	5	3%	3	2%	2	1%	1	1%
... None	104	72%	45	31%	86	59%	45	31%
Education/Veteran	145		145		145		145	
... HS Grad or Veteran	0	0%	2	1%	67	46%	105	72%
... 12th Grade or Veteran	0	0%	0	0%	2	1%	4	3%
... HS Grad	14	10%	17	12%	3	2%	5	3%
... HS Grad or Non-Veteran	0	0%	0	0%	4	3%	5	3%
... Other	5	3%	11	8%	8	6%	14	10%
... None	126	87%	115	79%	61	42%	12	8%
... HS Grad or In School	0	0%	0	0%	0	0%	0	0%
Years Continuous in USA	145		145		145		145	
... Used YRSUSA	10	7%	45	31%	18	12%	44	30%
... No YRSUSA	135	93%	100	69%	127	88%	101	70%

Multistep condition means the variable is one part of a complex boolean involving many different variables.

Table 13: Task 1 Effect and Samples by Sample Definitions

Variable	Treated-Group Restriction			All-Sample Restriction					
	Effect	Pctl. 25	Pctl. 50	Pctl. 75	Effect	Pctl. 25	Pctl. 50	Pctl. 75	Samp Size
Year of Immigration									
... < 2007	0.021		0.034	0.054	0.016		0.028	0.032	13,674
... <= 2007	0.013		0.028	0.052	0.012		0.019	0.036	37,376
Years Continuous in USA									
... Used YRSUSA	0.017		0.030	0.052	0.019		0.026	0.045	36,523
... No YRSUSA	0.012		0.030	0.047	0.014		0.030	0.051	69,522

Table 14: Task 2 Effect and Samples by Sample Definitions

Variable	Treated-Group Restriction			All-Sample Restriction					
	Effect	Pctl. 25	Pctl. 50	Pctl. 75	Effect	Pctl. 25	Pctl. 50	Pctl. 75	Samp Size
Year of Immigration									
... < 2007	0.015		0.028	0.048	0.014		0.029	0.053	22,840
... <= 2007	0.022		0.036	0.058	0.028		0.040	0.068	20,182
Years Continuous in USA									
... Used YRSUSA	0.008		0.034	0.058	0.025		0.034	0.063	15,424
... No YRSUSA	0.018		0.031	0.057	0.015		0.031	0.057	19,489

analysts style, with the two project organizers taking the same data and research question and performing independent analyses.

4.5.1 Analysis by Project Organizer A

For each categorical researcher characteristic specified in Section 3.2.5, as well as an indicator for the use of R or Stata as a programming language. In each case, categories with 5 or fewer researchers in them were omitted before performing the analysis. Table 15 shows the F-statistic from a regression of the reported effect estimate on a set of indicators for that characteristic, as well as the associated p -value and R^2 from that regression. This table allows us to see whether researchers with different characteristics reported different effect levels. Table 16 does the same, but uses absolute deviation from the sample mean as the dependent variable. This table allows us to see whether researchers with different characteristics showed greater agreement on effect levels with the group as a whole.

Table 15: Predicting Effect Level with Researcher Characteristics

Predictor	T1:			T2:			T3:		
	F	p	R2	F	p	R2	F	p	R2
Degree	0.929	0.337	0.007	0.122	0.727	0.001	0.085	0.771	0.001
Occupation	1.195	0.316	0.034	0.453	0.770	0.013	2.501	0.045	0.068
Research	1.080	0.342	0.015	0.370	0.692	0.005	0.416	0.660	0.006
Experience									
Gender	0.161	0.689	0.001	0.255	0.614	0.002	1.364	0.245	0.009

Predictor	T1:			T2:			T3:		
	F	p	R2	F	p	R2	F	p	R2
Race	1.026	0.383	0.022	1.306	0.275	0.028	0.342	0.795	0.007
LGBTQ+	0.426	0.654	0.006	0.183	0.833	0.003	0.045	0.956	0.001
Recruitment Source	0.360	0.698	0.005	1.661	0.194	0.024	1.400	0.250	0.020
Field	1.406	0.238	0.011	4.562	0.035	0.034	0.831	0.364	0.006
Coding Language	3.861	0.051	0.027	3.117	0.080	0.022	0.653	0.420	0.005

Table 16: Predicting Effect Deviation with Researcher Characteristics

Predictor	T1:			T2:			T3:		
	F	p	R2	F	p	R2	F	p	R2
Degree	1.915	0.169	0.013	0.740	0.391	0.005	0.630	0.429	0.004
Occupation	0.890	0.472	0.025	0.535	0.710	0.015	1.845	0.124	0.051
Research Experience	1.364	0.259	0.019	0.284	0.754	0.004	0.741	0.478	0.011
Gender	1.576	0.211	0.011	1.102	0.296	0.008	0.144	0.705	0.001
Race	2.180	0.093	0.045	0.129	0.943	0.003	0.762	0.517	0.016
LGBTQ+	0.202	0.817	0.003	0.515	0.599	0.007	0.253	0.776	0.004
Recruitment Source	2.064	0.131	0.030	0.197	0.822	0.003	0.552	0.577	0.008
Field	0.077	0.781	0.001	0.936	0.335	0.007	0.072	0.789	0.001
Coding Language	4.369	0.038	0.030	4.537	0.035	0.032	5.022	0.027	0.035

Tables 15 and 16 show that researcher characteristics hold basically no explanatory power for estimated effects either in level or deviation from the mean. Nearly all p -values are well above .05. In 16, the p -value for race as an explanatory variable in Task 1 had a p -value below .1, but given how many comparisons there are here, this is likely to just be noise.

The only researcher characteristic that did seem to matter was the choice of programming language, which only weakly predicted effect level, but was a statistically significant predictor of being close to the mean effect in all three rounds.

Figure 8 goes further into the split by language. We see that, of the two languages, Stata users were more likely to report effect estimates near the sample mean. 6.4%, 1.8%, and 0.9% of Stata users were more than .1 in absolute distance from the sample mean in Tasks 1, 2, and 3, respectively, while for R those values are 15.6%, 9.4%, and 12.5%. The number of R users is relatively low at 32,²¹ and so these numbers are sensitive to any researchers who were

²¹This is one lower than the value reported in Section 3.2.5 because the researcher who was dropped from analysis, mentioned later in Section 3.2.5, was an R user.




Figure 8: Deviation from Sample Mean of Reported Effect by Language

consistently outliers. There were two R users who had an absolute deviation from the mean of .1 or more every round, while all other R researchers with deviations of .1 or more only had deviations that large in a single round. If we omit those two consistently-high-deviation R users, the percentages are 10%, 3.3%, and 6.7% for R users, which are still higher than the percentages for Stata users.

Overall, there is little role for researcher professional or demographic characteristics in predicting either the level of the effects they reported, or the deviation of those effects from the mean. There is some explanatory power for the choice of programming language. R users were more likely than Stata users to report estimates far from average of what other users reported.

4.5.2 Analysis By Project Organizer B

Table 17 looks at within-researcher variation in effect estimates across tasks. In the first three columns, the dependent variable is the absolute difference in effects for a given researcher across two tasks, while in the fourth column, the dependent variable is a researcher's maximum estimated effect minus their minimum.

Most researcher characteristics do not predict absolute within-researcher variation. Career stage, occupation, and number of published papers do not predict absolute differences in estimates across tasks to a statistically significant degree, with few exceptions.

One exception is that private researchers saw larger absolute changes between Task 1 and Task 3, and also more absolute variation overall, although the latter is only significant at the $\alpha = .1$ level. Probably the most interesting is that inexperience was related to smaller changes from Task 1 to Task 3: those who do not have a PhD showed a smaller change between Task 1 and Task 3 (significant at $\alpha = .1$), and those with fewer papers also showed smaller absolute changes than those with 6+ papers (insignificant).

4.5.3 Two-Analyst Results for Researcher Characteristics

Both project organizers used the same data to answer the same research question. From the preregistration: "Both primary authors will, independently, analyze the relationship between (a) researcher characteristics and reported research results in earlier stages, and (b) attrition from the study and reported research results in later stages." Because there was so little attrition from the study after Task 1, part b was dropped from the analysis. While this question is much more open-ended than the main causal inference question in this paper, it is notable how different the two approaches were. The dependent variable of interest differed, as did the analytic method and selection of relevant predictors. Both organizers found, however, that researcher characteristics were not strong predictors of what the researchers did.

Table 17: Model Coefficients and Standard Errors for Task Comparisons

	Task 1 vs Task 2	Task 2 vs Task 3	Task 1 vs Task 3	Absolute Range
Intercept	0.053*** (0.015)	0.068*** (0.018)	0.053** (0.017)	0.087*** (0.022)
Grad student	0.090+ (0.047)	−0.015 (0.054)	0.099+ (0.051)	0.086 (0.066)
Uni researcher	−0.001 (0.033)	−0.014 (0.038)	0.016 (0.036)	0.000 (0.047)
Other	0.004 (0.070)	−0.013 (0.081)	0.032 (0.077)	0.011 (0.099)
Private researcher	0.021 (0.042)	0.065 (0.049)	0.134** (0.046)	0.110+ (0.059)
Public researcher	0.047 (0.035)	−0.013 (0.041)	0.044 (0.039)	0.039 (0.050)
Not PhD	−0.070+ (0.040)	0.003 (0.047)	−0.081+ (0.044)	−0.074 (0.057)
1-5 papers	−0.007 (0.021)	−0.037 (0.025)	−0.011 (0.024)	−0.027 (0.030)
0 papers	0.012 (0.032)	−0.047 (0.037)	−0.017 (0.035)	−0.026 (0.045)
Num.Obs.	145	145	145	145
R2	0.046	0.040	0.088	0.047
R2 Adj.	−0.010	−0.017	0.034	−0.009

+ $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

4.6 Bimodality in the Task 2 Effect Estimates

One of the surprising results in Section 4.1 was the effect distribution in Task 2. In designing the study, we had expected that each task would show a narrower distribution of effects than the previous task. While we did generally see this pattern for sample sizes and some researcher choices, the distribution of effects in particular became wider going from Task 1 to Task 2. We also saw some emerging bimodality, where the larger part of the sample reported estimates that reflected the distribution of effects already seen in Task 1, while a smaller group of researchers reported larger effects that were more like those found in Task 3. In this section we explore possible explanation for the unexpected findings in Task 2.²²

The fact that the effect distribution is not just wider but rather gathers at a high and a low point makes the task of explaining it somewhat easier, as we can look for features that predict reporting a higher or lower estimate.

Several anticipated correlates did not explain the bimodal outcomes of Task 2. Figure 11 and Figure 12 in the Appendix show that the Task 2 reported sample sizes and standard errors do not strongly explain the effects reported. Recalling Section 4.3, there were not enormous differences in reported estimates by control variables included, and there were not covariates associated with large effects that increased in prominence in Task 2, so it is unlikely that changes in covariate choice or importance explain the effect distribution.

A possible expansion of the bimodality is that some researchers found their Task 1 analyses or results “sticky” and tried to match them too closely, while researchers who did not attempt to do this instead produced results like those in Task 3, since the Task 2 instructions are very similar to Task 3. However, Figure 9 shows effectively no relationship between a given researcher’s Task 1 estimate and their Task 2 estimate.

We then look at sample limitations and definitions. Table 18 uses the sample-restriction coding from Section 4.4, and examines whether a given researcher changed their use of a given variable in their sample restrictions from Task 1 to Task 2. This could be any change, for example going from not using the “Hispanic” variable at all to using it, or switching from “Hispanic-Mexican” to “Hispanic-Any”. The table shows that, with Hispanic and Birthplace as exceptions, the share above .05 tended to be higher among researchers who changed the way they used a given variable in defining their sample.

Table 18: Changes in Sample Limitations from Task 1 to 2, and Effect Changes

Sample Limitation	Changed	N	Increase	Standard Error	R2Effect	Abovep05
Hispanic	No	110	-0.01	0.01	0.04	36.4%
Hispanic	Yes	35	0.00	0.02	0.04	22.9%
Birthplace	No	112	0.00	0.01	0.05	35.7%

²²This section is entirely un-preregistered, as we did not anticipate this finding.

Sample Limitation	Changed	N	Increase	Standard Error	R2Effect	Abovep05
Birthplace	Yes	33	-0.03	0.01	0.03	24.2%
Citizenship	No	106	0.00	0.01	0.05	31.1%
Citizenship	Yes	39	-0.02	0.02	0.04	38.5%
Age at Migration	No	75	-0.01	0.01	0.04	29.3%
Age at Migration	Yes	70	-0.01	0.02	0.05	37.1%
Age in June 2012	No	71	-0.01	0.01	0.04	31.0%
Age in June 2012	Yes	74	-0.01	0.02	0.05	35.1%
Year of Immigration	No	88	-0.01	0.01	0.04	28.4%
Year of Immigration	Yes	57	0.00	0.02	0.05	40.4%
Education/Veteran	No	57	0.00	0.01	0.06	31.6%
Education/Veteran	Yes	88	-0.02	0.02	0.03	34.1%
Years Continuous in USA	No	129	-0.01	0.01	0.05	32.6%
Years Continuous in USA	Yes	16	0.00	0.02	0.04	37.5%

Taking the explanatory power of sample restrictions, we then look at the treated-group definition. Task 2 gave a very precise definition of who should be included as a part of the treated group. We examine whether a given researcher followed the full set of treated-group definition instructions precisely or not. The mismatch could be small, such as using “ ≤ 16 ” instead of “ < 16 ” for age at migration, or large, such as omitting that eligible people must be non-citizens. In Figure 10 we show their distribution of effects against researchers who had a mismatch in their criteria in any way. The graph shows that the bimodality heavily driven by the group that precisely matched the treated-group definition. This implies that the bimodality in Task 2 may be explained in large part by a split between researchers who exactly followed the instructions, and so effectively matched what a typical researcher found in Task 3, and those who did not. This does not fully explain researcher behavior: note that there is also a weight of researchers with higher results who did not match perfectly, and also that much of the density of the perfect-match group is at Task 1 effect levels, but keep in mind that there are many other decisions in analysis to be made, and this captures only one angle where determining the correct decision is easiest.

Further, Table 19 shows that the share of researchers matching exactly is fairly low, between 20-25% by field, keeping in mind that even very minor mismatches are counted as mismatches. Further, perfect-match rates were slightly higher among researchers whose work was closest to the field that the research task was in, immigration and labor, although this difference was not statistically significant at the 95% level.

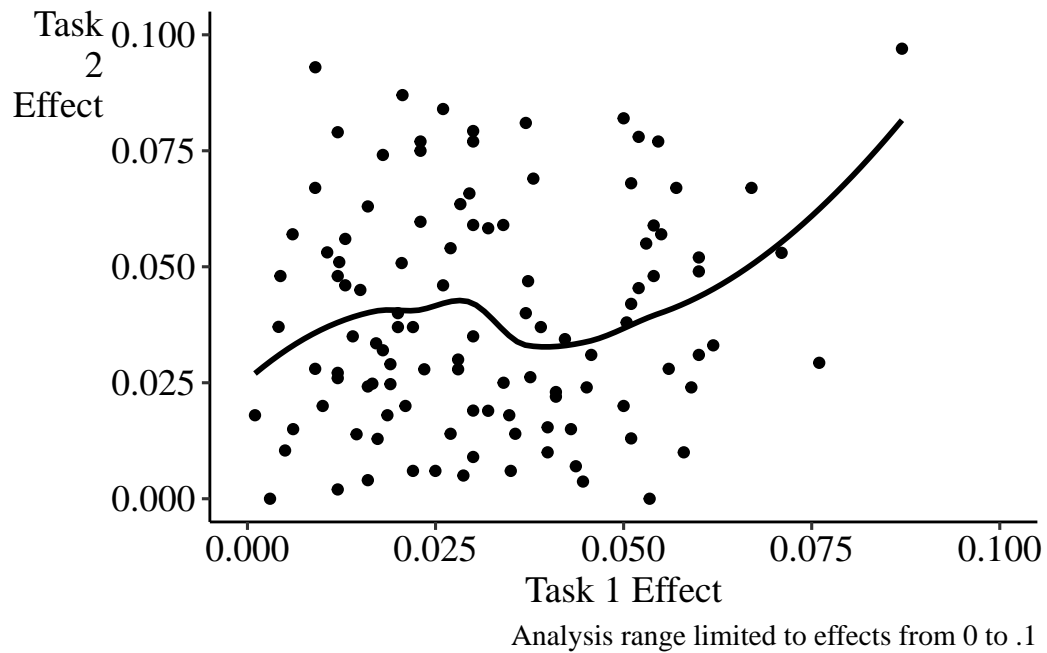


Figure 9: Estimates Effects in Task 1 vs. Task 2

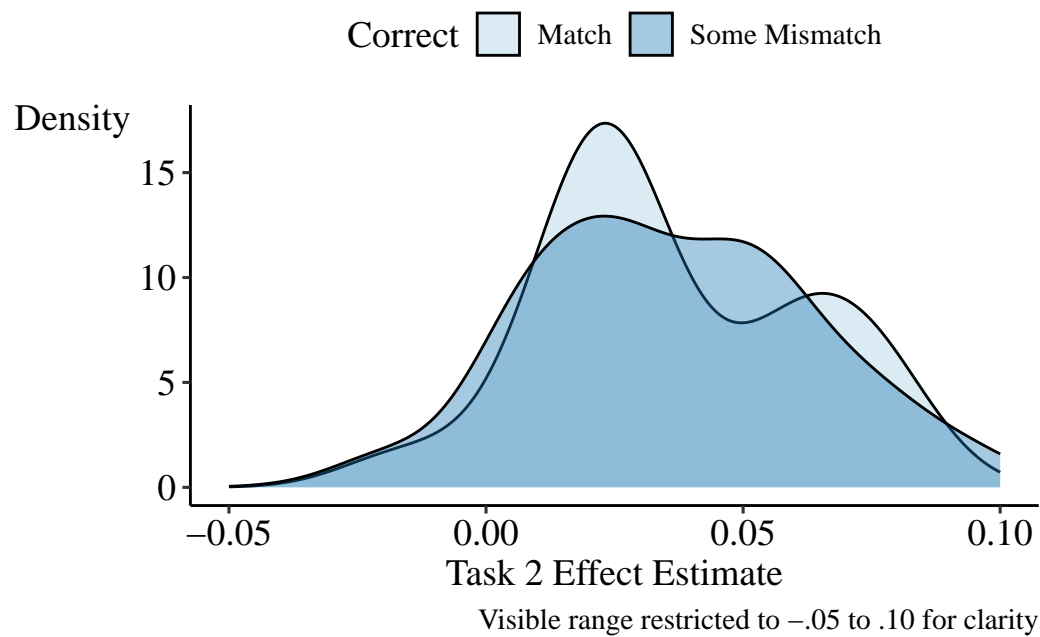


Figure 10: Task 2 Effect Distributions Among Those with Exact Treated-Group Definition Matches vs. Those with Some Mismatch

Table 19: Share of Researchers Matching Treated-Group Definition Exactly by Field

Field	Share Match	Num. Match	Share Some Mismatch	Num Some Mismatch
Immigration	25.0%	1	75.0%	3
Labor	25.5%	12	74.5%	35
Neither/Other	22.3%	21	77.7%	73

5 Conclusion

5.1 Recommendations for Improved Practice

What do these results imply should change about the practice of applied microeconomic research?

To some degree, the findings of this paper do not reflect a problem to be solved. The fact that different researchers approach a problem differently is not in itself a problem, as long as any points of disagreement are visible to the reader and subject to scrutiny and disagreement, and the reader understands that a given study or set of research decisions is not the last word.

However, there is a problem to be solved to the extent that researcher variation reflects either (a) error, or (b) choices that are unexamined or invisible while also being something that researchers would choose differently.

In this study, we found considerable variation across researchers in the approaches taken to answering the main question in Task 1, which most closely reflects actual practice. While the distribution of effects did not considerably narrow from Task 1 to Task 2, the sample sizes did, as did the treated and comparison group definitions. Although we generally did not reject Levene tests of equal variance across rounds, descriptively there was an obvious narrowing of the distribution of effects. There was, in both rounds, a lot of variation in the set of covariates included, as well.

To the extent that these choices reflect research design and modeling choices, the problem is somewhat already addressed. Researchers are used to critiquing research design and modeling choices in public work. Because of this, we would expect that all of these choices would be reported in a writeup of research, where they could be critiqued. What would not typically occur is someone actually testing whether many of these alternate choices actually lead to different results, especially for seemingly more innocuous choices like covariate functional form, which was found in this study to be more consequential than the set of covariates included.

On the part of researcher practice, this set of results suggests the use of multiverse analysis (Steege et al. 2016), where the researcher considers every combination of reasonable modeling decisions and demonstrates their effects on estimated effects, or even many-analysts approaches

to producing original work, as we did in Section Section 4.5. On the part of journals, this suggests that journals should consider accepting work that is a variation in the approach to a published work, even if that variation is not framed as a replication or rejection of the original study, currently a barrier to the publication of replications (Galiani, Gertler, and Romero 2017).

However, this study did not find that research design and modeling choices were the majority of explained researcher variation. Instead, this came in the form of data cleaning and preprocessing, including the selection a sample and the creation of variables indicating the treated group. Some of this variation could be classified as error, for example researchers in Section Section 4.6 whose treated-group definition in Task 2 did not match the instructions. Other parts of this variation could be reasonable disagreement.

That much of the relevant variation seems to come along the lines of data cleaning and preprocessing is possibly unsurprising, given how this task is currently handled in economics. Relative to modeling and research design, data cleaning and preprocessing receive little attention. It is common for minor, or even important, details of data processing to be left of the description of methods in published papers. Data cleaning is also not formally taught in most graduate programs. A professor who would never allow a research assistant to decide their research design or model might be happy passing along the data cleaning task to an assistant, even though, as this paper shows, the task may be just as relevant to the results and just as prone to arbitrary choice-making.

Because data cleaning gets so little attention, it is perhaps to be expected that we saw the most variation here. Without formal training in PhD programs, or a culture of reviewing and critiquing data cleaning and preprocessing in research papers, there is little opportunity for researchers to *learn to do the same thing*, and so we see heavy variation. Contrast this to the popular use of linear probability models in Table 7, for example, which is common in applied microeconomics, especially in difference-in-differences designs. Because its use is very visible, researchers can see that it is the standard in the field. Whether or not it is actually the best method to use here, it is a method that researchers know has been agreed upon in the field. We see a similar story play out with standard error adjustments in the same table. The standard method by which economists learn data processing is on a much narrower scale, often from one’s advisor or from others on a small research team.

This problem implies several possible policy solutions. Among broader system-wide changes, the introduction of data cleaning and preprocessing classes teaching methods and best practices in the standard PhD applied economics curriculum would likely improve the quality of economics research as well as reduce researcher variability. This presumes the existence of a set of best practices, or at least standard practices. So, an attempt should be made to codify and popularize a set of data-cleaning best practices, similar to how applied economists routinely learn about modeling best practices (for example any number of econometrics textbooks, or in the applied literature papers like Abadie et al. (2023)). Other fields have already made strides in this direction (e.g. Osborne 2012; Jafari 2022) so this effort would not need to start from

scratch, but would be improved by designing recommendations most relevant to an applied microeconomics context.

Those recommendations may not be in the control of any one researcher, though. The policy that this paper implies for individual researchers and journal editors or reviewers is that economics as a field should consider data cleaning and preprocessing to be just as much a part of the methods as the choice of model. Researchers should fully describe their data cleaning processes in their papers to the same level of detail that they describe their modeling choices. They should also perhaps subject any arbitrary data-cleaning decisions to multiverse analysis as well (note that this is also a suggestion of Steegen et al. (2016)). Further, an increasingly common requirement of journal submissions is to provide a replication package of the code used to perform a paper’s analyses (for example American Economic Association (2024)). However, these replication packages often begin from an already-prepared data set, and only include the code necessary to run models. It would be advisable to include data preprocessing code in these replication packages.

5.2 Discussion

This paper describes the results of a large many-analysts project in applied microeconomics. In this project we found large amounts of variation in the choices made by researchers, especially in regards to data cleaning and processing, research design, the definition of treated and comparison groups, and the selection and functional form of controls. Some of this variation appears to be from researcher data cleaning processes that do not match the instructions, derived from policy realities, for constructing the treated group. Variation was not strongly constrained by the influence of peer review or a shared research design, but there was a (descriptive if not statistically significant) reduction of variation when researchers were provided with pre-cleaned data.

Interestingly, we do not find huge amounts of variation in actual estimated effects of the policy. There were some outliers, and statistical significance varied between researchers. However, the central range of estimated effects of DACA on the probability of working full time effects generally was not very wide, with the difference between the 25th and 75th percentiles typically only 2-3 percentage points. However, the fact that widely different sample definitions and modeling choices led to a narrow range of effects is not guaranteed to generalize to other contexts.

There were also parts where researchers behaved very similarly. The use of linear regression modeling was very popular, and very few researchers used unadjusted standard errors, although the specific adjustment or clustering level varied.

What we might understand from this study, and perhaps the wider world of studies on researcher variation, is perhaps obvious: in cases where there were well-acknowledged “standard” ways of doing something, like using linear modeling in a difference-in-differences-type setting with a binary outcome, or adjusting one’s standard errors, researchers tended to do that thing.

And where there was no well-acknowledged standard, like in the choice of clustering level, the selection of covariates in this particular setting, or in data cleaning, researchers behaved differently, sometimes to consequence and sometimes without it mattering much.

While researcher error also plays a part, the impact of the absence of standards is an emergent result from this study. Readers and practitioners of research should expect arbitrary variation in parts of research, like data cleaning, that do not have standards.

The development of best-practice standards in areas where we currently do not acknowledge them would be likely to improve applied microeconomics towards being a more mature, sophisticated, and believable field than it is today. We have highlighted data-cleaning practices as being an especially fruitful place to develop these standards, but the same applies in other areas as well like the level of clustering (an example of a place where development of consensus guidance is already underway in Abadie et al. (2023)). The optimal level of researcher variation is not zero, as individual researchers often have good reasons not to match the methods and practices used by others. But when this occurs, it should be because there *is* a good reason to deviate from the template, rather than because we have no template to begin with.

References

- Abadie, Alberto, Susan Athey, Guido W Imbens, and Jeffrey M Wooldridge. 2023. “When Should You Adjust Standard Errors for Clustering?” *The Quarterly Journal of Economics* 138 (1): 1–35.
- American Economic Association. 2024. “Data and Code Availability Policy.” AEA. <https://www.aeaweb.org/journals/data/data-code-policy>.
- Amuedo-Dorantes, Catalina, and Francisca Antman. 2016. “Can Authorization Reduce Poverty Among Undocumented Immigrants? Evidence from the Deferred Action for Childhood Arrivals Program.” *Economics Letters* 147: 1–4.
- Ankel-Peters, Jörg, Nathan Fiala, and Florian Neubauer. 2023. “Do Economists Replicate?” *Journal of Economic Behavior & Organization* 212: 219–32.
- Arel-Bundock, Vincent. 2022. “modelssummary: Data and Model Summaries in R.” *Journal of Statistical Software* 103 (1): 1–23. <https://doi.org/10.18637/jss.v103.i01>.
- Auspurg, Katrin, and Josef Brüderl. 2021. “Has the Credibility of the Social Sciences Been Credibly Destroyed? Reanalyzing the ”Many Analysts, One Data Set” Project.” *Socius* 7: 23780231211024421.
- . 2023. “Is Social Research Really Not Better Than Alchemy? How Many-Analysts Studies Produce ‘a Hidden Universe of Uncertainty’ by Not Following Meta-Analytical Standards.”
- Barrett, Tyson, Matt Dowle, Arun Srinivasan, Jan Gorecki, Michael Chirico, and Toby Hocking. 2024. *Data.table: Extension of data.frame*. <https://r-datatable.com>.
- Bastiaansen, Jojanneke A, Yoram K Kunkels, Frank J Blaauw, Steven M Boker, Eva Ceulemans, Meng Chen, Sy-Miin Chow, et al. 2020. “Time to Get Personal? The Impact of

- Researchers Choices on the Selection of Treatment Targets Using the Experience Sampling Methodology.” *Journal of Psychosomatic Research* 137: 110211.
- Becker, Jason, Chung-hong Chan, David Schoch, and Thomas J. Leeper. 2023. *Rio: A Swiss-Army Knife for Data i/o*. <https://github.com/gesistsa/rio>.
- Bergé, Laurent. 2018. “Efficient Estimation of Maximum Likelihood Models with Multiple Fixed-Effects: The R Package FENmlm.” *CREA Discussion Papers*, no. 13.
- Boehm, Udo, Jeffrey Annis, Michael J Frank, Guy E Hawkins, Andrew Heathcote, David Kellen, Angelos-Miltiadis Kryptos, et al. 2018. “Estimating Across-Trial Variability Parameters of the Diffusion Decision Model: Expert Advice and Recommendations.” *Journal of Mathematical Psychology* 87: 46–75.
- Botvinik-Nezer, Rotem, Felix Holzmeister, Colin F Camerer, Anna Dreber, Juergen Huber, Magnus Johannesson, Michael Kirchler, et al. 2020. “Variability in the Analysis of a Single Neuroimaging Dataset by Many Teams.” *Nature* 582 (7810): 84–88.
- Breznau, Nate, Eike Mark Rinke, Alexander Wuttke, Muna Adem, Jule Adriaans, Amalia Alvarez-Benjumea, Henrik Kenneth Andersen, et al. 2021. “Observing Many Researchers Using the Same Data and Hypothesis Reveals a Hidden Universe of Data Analysis.” *MetaArXiv Preprints*.
- Broderick, Tamara, Ryan Giordano, and Rachael Meager. 2020. “An Automatic Finite-Sample Robustness Metric: When Can Dropping a Little Data Make a Big Difference?” *arXiv Preprint arXiv:2011.14999*.
- Bryan, Christopher J, David S Yeager, and Joseph M O’Brien. 2019. “Replicator Degrees of Freedom Allow Publication of Misleading Failures to Replicate.” *Proceedings of the National Academy of Sciences* 116 (51): 25535–45.
- Callaway, Brantly, and Pedro HC Sant’Anna. 2021. “Difference-in-Differences with Multiple Time Periods.” *Journal of Econometrics* 225 (2): 200–230.
- Camerer, Colin F, Anna Dreber, Eskil Forsell, Teck-Hua Ho, Jürgen Huber, Magnus Johannesson, Michael Kirchler, et al. 2016. “Evaluating Replicability of Laboratory Experiments in Economics.” *Science* 351 (6280): 1433–36.
- Card, David, Stefano DellaVigna, Patricia Funk, and Nagore Iriberri. 2022. “Gender Differences in Peer Recognition by Economists.” *Econometrica* 90 (5): 1937–71.
- Chen, Wanyi, and Mary Cummings. 2024. “Subjectivity in Unsupervised Machine Learning Model Selection.” In *Proceedings of the AAAI Symposium Series*, 3:22–29. 1.
- Fox, John, and Sanford Weisberg. 2019. *An R Companion to Applied Regression*. Third. Thousand Oaks CA: Sage. <https://socialsciences.mcmaster.ca/jfox/Books/Companion/>.
- Galiani, Sebastian, Paul Gertler, and Mauricio Romero. 2017. “Incentives for Replication in Economics.” National Bureau of Economic Research.
- Gould, Elliot, Hannah S Fraser, Timothy H Parker, Shinichi Nakagawa, Simon C Griffith, Peter A Vesk, Fiona Fidler, et al. 2023. “Same Data, Different Analysts: Variation in Effect Sizes Due to Analytical Decisions in Ecology and Evolutionary Biology.”
- Herbert, Sylvérie, Hautahi Kingi, Flavio Stanchi, and Lars Vilhuber. 2021. “The Reproducibility of Economics Research: A Case Study.”
- Holzmeister, Felix, Magnus Johannesson, Robert Böhm, Anna Dreber, Jürgen Huber, and Michael Kirchler. 2023. “Heterogeneity in Effect Size Estimates: Empirical Evidence and

- Practical Implications.” Working Papers in Economics; Statistics.
- Hoogeveen, Suzanne, Alexandra Sarafoglou, Balazs Aczel, Yonathan Aditya, Alexandra J Alayan, Peter J Allen, Sacha Altay, et al. 2023. “A Many-Analysts Approach to the Relation Between Religiosity and Well-Being.” *Religion, Brain & Behavior* 13 (3): 237–83.
- Huntington-Klein, Nick. 2021. *Vtable: Variable Table for Variable Documentation*. <https://nickch-k.github.io/vtable/>.
- Huntington-Klein, Nick, Andreu Arenas, Emily Beam, Marco Bertoni, Jeffrey R Bloem, Pralhad Burli, Naibin Chen, et al. 2021. “The Influence of Hidden Researcher Decisions in Applied Microeconomics.” *Economic Inquiry* 59 (3): 944–60.
- Jafari, Roy. 2022. *Hands-on Data Preprocessing in Python: Learn How to Effectively Prepare Data for Successful Data Analytics*. Packt Publishing Ltd.
- Jelveh, Zubin, Bruce Kogut, and Suresh Naidu. 2024. “Political Language in Economics.” *The Economic Journal*, ueae026.
- Klau, Simon, Chirag J Patel, John PA Ioannidis, Anne-Laure Boulesteix, Sabine Hoffmann, et al. 2023. “Comparing the Vibration of Effects Due to Model, Data Pre-Processing and Sampling Uncertainty on a Large Data Set in Personality Psychology.” *Meta-Psychology* 7.
- Leamer, Edward E. 1983. “Let’s Take the Con Out of Econometrics.” *The American Economic Review* 73 (1): 31–43.
- Lundberg, Shelly, and Jenna Stearns. 2019. “Women in Economics: Stalled Progress.” *Journal of Economic Perspectives* 33 (1): 3–22.
- Menkveld, Albert J, Anna Dreber, Felix Holzmeister, Juergen Huber, Magnus Johannesson, Michael Kirchler, Sebastian Neusüss, Michael Razen, and Utz Weitzel. 2021. “Non-Standard Errors.”
- Open Science Collaboration. 2015. “Estimating the Reproducibility of Psychological Science.” *Science* 349 (6251): aac4716.
- Ortloff, Anna-Marie, Matthias Fassl, Alexander Ponticello, Florin Martius, Anne Mertens, Katharina Krombholz, and Matthew Smith. 2023. “Different Researchers, Different Results? Analyzing the Influence of Researcher Experience and Data Type During Qualitative Analysis of an Interview and Survey Study on Security Advice.” In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*, 1–21.
- Osborne, Jason W. 2012. *Best Practices in Data Cleaning: A Complete Guide to Everything You Need to Do Before and After Collecting Your Data*. Sage Publications.
- Ostropolets, Anna, Yasser Albogami, Mitchell Conover, Juan M Banda, William A Baumgartner Jr, Clair Blacketer, Priyamvada Desai, et al. 2023. “Reproducible Variability: Assessing Investigator Discordance Across 9 Research Teams Attempting to Reproduce the Same Observational Study.” *Journal of the American Medical Informatics Association* 30 (5): 859–68.
- Pérignon, Christophe, Olivier Akmansoy, Christophe Hurlin, Anna Dreber, Felix Holzmeister, Juergen Huber, Michael Kirchler, Michael Razen, et al. 2022. “Reproducibility of Empirical Results: Evidence from 1,000 Tests in Finance.”
- Portner, Claus C, and Nick Huntington-Klein. 2022. “Many Economists.” OSF. <https://doi.org/10.31233/osf.io/zt4qj>.

[org/10.17605/OSF.IO/CJ9YX](https://osf.io/CJ9YX).

- Ruggles, Steven, Sarah Flood, Matthew Sobek, Daniel Backman, Annie Chen, Grace Cooper, Stephanie Richards, Renae Rodgers, and Megan Schouweiler. 2024. "IPUMS USA: Version 15.0 [Dataset]. Minneapolis, MN: IPUMS."
- Sarafoglou, Alexandra, Suzanne Hoogeveen, Don Van Den Bergh, Balázs Aczél, Casper J Albers, Tim Althoff, Rotem Botvinik-Nezer, et al. 2024. "Subjective Evidence Evaluation Survey for Multi-Analyst Studies."
- Schweinsberg, Martin, Michael Feldman, Nicola Staub, Olmo R van den Akker, Robbie CM van Aert, Marcel ALM Van Assen, Yang Liu, et al. 2021. "Same Data, Different Conclusions: Radical Dispersion in Empirical Results When Independent Analysts Operationalize and Test the Same Hypothesis." *Organizational Behavior and Human Decision Processes* 165: 228–49.
- Silberzahn, Raphael, Eric L Uhlmann, Daniel P Martin, Pasquale Anselmi, Frederik Aust, Eli Awtrey, Štěpán Bahník, et al. 2018. "Many Analysts, One Data Set: Making Transparent How Variations in Analytic Choices Affect Results." *Advances in Methods and Practices in Psychological Science* 1 (3): 337–56.
- Simmons, Joseph P, Leif D Nelson, and Uri Simonsohn. 2011. "False-Positive Psychology: Undisclosed Flexibility in Data Collection and Analysis Allows Presenting Anything as Significant." *Psychological Science* 22 (11): 1359–66.
- Stansbury, Anna, and Robert Schultz. 2023. "The Economics Profession's Socioeconomic Diversity Problem." *Journal of Economic Perspectives* 37 (4): 207–30.
- Steege, Sara, Francis Tuerlinckx, Andrew Gelman, and Wolf Vanpaemel. 2016. "Increasing Transparency Through a Multiverse Analysis." *Perspectives on Psychological Science* 11 (5): 702–12.
- Sulik, Justin, Nakwon Rim, Elizabeth Pontikes, James Evans, and Gary Lupyan. 2023. "Why Do Scientists Disagree?"
- Urban Institute. 2022. "State Immigration Policy Resource." <https://www.urban.org/data-tools/state-immigration-policy-resource>.
- U.S. Citizenship and Immigration Services. 2016. "Number of i-821D,consideration of Deferred Action for Childhood Arrivals by Fiscal Year, Quarter, Intake, Biometrics and Case Status: 2012-2016 (June 30)."
- Wickham, Hadley, Mara Averick, Jennifer Bryan, Winston Chang, Lucy D'Agostino McGowan, Romain François, Garrett Grolemund, et al. 2019. "Welcome to the tidyverse." *Journal of Open Source Software* 4 (43): 1686. <https://doi.org/10.21105/joss.01686>.

Appendix

Table 20: Paired Absolute Effect Differences and Peer Review

	Task 1	Task 2
Intercept	0.086*** (0.010)	0.063*** (0.009)
Comparison: Next Round	-0.034** (0.014)	-0.006 (0.013)
Comparison: Next Round vs. This Round	-0.017 (0.014)	0.000 (0.013)
Unreviewed	-0.045*** (0.010)	0.001 (0.010)
Next Round x Unreviewed	0.057*** (0.014)	-0.030** (0.014)
Next vs. This x Unreviewed	0.029** (0.014)	-0.017 (0.014)
Num.Obs.	6809	6676

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$




Figure 11: Task 2 Effect Size and Sample Size

Table 21: Task 1 Effect and Samples by Sample Definitions, Full View

Variable	Treated-Group Restriction			All-Sample Restriction			Samp Size	Pctl. 25	Pctl. 50	Pctl. 75
	Effect	Pctl. 25	Pctl. 50	Effect	Pctl. 25	Pctl. 50				
Hispanic										
... Hispanic-Mexican	0.013		0.030	0.053	0.015	0.032	0.056	64,614	158,152	292,492
... Hispanic-Any	0.030		0.037	0.050	0.030	0.045	0.054	24,276	109,759	166,270
... Hispanic-Mex or Mex-Born	0.018		0.021	0.026	0.021	0.022	0.027	51,754	127,504	179,960
... Multistep Condition	-0.009		0.001	0.010	-0.009	0.001	0.010	326,913	366,804	406,696
... None	0.015		0.023	0.052	0.012	0.022	0.048	65,124	259,630	767,657
Birthplace										
... Mexican-Born	0.016		0.029	0.046	0.016	0.030	0.050	61,412	141,847	276,683
... Hispanic-Mex or Mex-Born	0.041		0.100	0.160	0.004	0.016	0.070	215,436	289,311	330,348
... Non-US Born	-0.020		-0.020	-0.020	-0.013	-0.007	-0.001	133,992	157,710	181,428
... Central America-Born	0.057		0.057	0.057	0.057	0.057	0.057	9,711	9,711	9,711
... None	0.012		0.032	0.053	0.011	0.034	0.052	70,156	277,277	746,663
Citizenship										
... Non-Citizen	0.017		0.030	0.051	0.022	0.035	0.056	62,920	141,847	277,270
... Foreign-Born	0.026		0.037	0.180	0.026	0.037	0.180	341,338	586,271	605,241
... Non-Cit or Natlzd post-2012	0.001		0.022	0.028	0.017	0.017	0.017	13,377	13,377	13,377
... Citizen (various)	0.019		0.023	0.047	0.015	0.015	0.015	268,238	268,238	268,238
... Multistep Condition	0.009		0.009	0.009	-0.034	-0.020	-0.005	899,372	1,694,116	2,488,861
... Other	0.023		0.041	0.053	0.023	0.037	0.046	13,427	17,759	84,944
... None	0.010		0.021	0.051	0.010	0.020	0.043	83,476	225,894	656,161
Age at Migration										
... < 16	0.016		0.030	0.050	0.018	0.026	0.046	28,348	44,288	119,180
... <= 16	0.010		0.022	0.044	0.034	0.041	0.044	120,931	204,239	205,147
... Other	0.019		0.035	0.058	0.012	0.028	0.054	40,548	117,257	208,712
... None	0.002		0.029	0.050	0.013	0.030	0.052	95,336	255,769	507,856
Age in June 2012										
... Year-Quarter Age	0.015		0.030	0.052	0.017	0.029	0.051	40,661	132,637	255,734
... Year-Only Age	0.010		0.030	0.041	0.013	0.037	0.053	17,759	46,925	140,134
... None	0.020		0.029	0.048	0.015	0.030	0.051	90,330	206,266	424,859
Year of Immigration										
... < 2007	0.021		0.034	0.054	0.016	0.028	0.032	13,674	32,242	51,475
... <= 2007	0.013		0.028	0.052	0.012	0.019	0.036	37,376	82,717	205,986
... < 2012	0.014		0.014	0.014	0.039	0.051	0.064	88,982	103,534	118,086
... <= 2012	0.042		0.118	0.194	0.029	0.092	0.155	263,220	471,364	679,507
... >= 2007	0.012		0.012	0.012	0.012	0.012	0.012	245,635	245,635	245,635
... Any Year	0.018		0.030	0.040	0.021	0.034	0.044	31,170	116,405	212,998
... Multistep Condition	0.014		0.014	0.014	0.010	0.012	0.013	137,786	170,943	204,100
... Other	-0.004		0.010	0.175	0.010	0.020	0.034	61,225	139,544	338,618
... None	0.017		0.030	0.050	0.016	0.032	0.053	110,144	230,665	474,472
Education/Veteran										
... HS Grad or Veteran	0.043		0.065	0.088						
... HS Grad	0.014		0.036	0.050	0.016	0.037	0.052	66,766	132,990	169,327
... Other	0.020		0.027	0.059	0.027	0.057	0.062	32,606	74,431	188,802
... None	0.014		0.030	0.051	0.014	0.030	0.051	62,354	203,345	387,252
Years Continuous in USA										
... Used YRSUSA	0.017		0.030	0.052	0.019	0.026	0.045	36,523	83,097	151,054
... No YRSUSA	0.012		0.030	0.047	0.014	0.030	0.051	69,522	194,349	367,941

Table 22: Task 2 Effect and Samples by Sample Definitions, Full View

Variable	Treated-Group Restriction					All-Sample Restriction						
	Effect	Pctl. 25	Pctl. 50	Pctl. 75	Effect	Pctl. 25	Pctl. 50	Pctl. 75	Samp Size	Pctl. 25	Pctl. 50	Pctl. 75
Hispanic												
... Hispanic-Mexican	0.015		0.029	0.057	0.014		0.029	0.057	18,981		25,199	43,220
... Hispanic-Any	0.028		0.050	0.063	0.024		0.042	0.058	22,983		24,011	25,568
... Hispanic-Mex or Mex-Born	0.048		0.048	0.048								
... Multistep Condition	0.025		0.025	0.025	0.025		0.025	0.025	44,805		44,805	44,805
... None	0.020		0.040	0.067	0.023		0.043	0.075	19,106		26,396	63,634
Birthplace												
... Mexican-Born	0.018		0.032	0.056	0.018		0.030	0.056	19,402		25,155	38,154
... Hispanic-Mex or Mex-Born	0.070		0.070	0.070								
... Non-US Born				0.028		0.045	0.058		25,498		26,138	47,180
... Central America-Born	0.067		0.067	0.067	0.067		0.067	0.067	25,538		25,538	25,538
... None	0.014		0.030	0.058	0.014		0.035	0.059	18,680		27,152	75,755
Citizenship												
... Non-Citizen	0.018		0.031	0.058	0.019		0.033	0.058	18,834		24,979	41,917
... Foreign-Born				0.004		0.004	0.004		164,135		164,135	164,135
... Non-Cit or Natlzd post-2012	0.014		0.034	0.048	0.019		0.024	0.029	18,162		21,823	25,484
... Citizen (various)	0.045		0.045	0.045	0.045		0.045	0.045	19,168		19,168	19,168
... Other	0.020		0.049	0.059	0.031		0.059	0.059	21,828		24,011	90,368
... None	0.007		0.021	0.045	0.013		0.028	0.053	19,855		28,756	75,402
Age at Migration												
... < 16	0.017		0.032	0.057	0.027		0.041	0.059	19,402		23,350	27,536
... <= 16	0.008		0.021	0.045	0.014		0.020	0.038	20,392		25,199	26,267
... Other	0.019		0.039	0.072	0.023		0.036	0.055	16,365		24,249	49,067
... None	0.023		0.037	0.065	0.013		0.026	0.053	19,367		31,458	94,773
Age in June 2012												
... Year-Quarter Age	0.019		0.034	0.058	0.024		0.036	0.057	19,864		25,424	42,427
... Year-Only Age	-0.002		0.018	0.042	0.015		0.019	0.048	16,542		26,396	29,146
... None	0.008		0.018	0.049	0.007		0.026	0.059	18,803		25,414	86,316
Year of Immigration												
... < 2007	0.015		0.028	0.048	0.014		0.029	0.053	22,840		25,056	28,602
... <= 2007	0.022		0.036	0.058	0.028		0.040	0.068	20,182		25,588	28,895
... < 2012	0.038		0.047	0.055								
... <= 2012	-0.011		0.068	0.290	0.068		0.068	0.068	6,600		6,600	6,600
... >= 2007	0.014		0.014	0.014								
... Any Year	0.026		0.037	0.048	0.049		0.052	0.056	18,409		20,276	22,144
... Multistep Condition	0.000		0.017	0.035								
... Other	0.850		0.850	0.850	0.024		0.026	0.028	25,048		35,419	45,790
... None	0.013		0.035	0.057	0.013		0.028	0.056	18,803		26,127	72,235
Education/Veteran												
... HS Grad or Veteran	0.018		0.031	0.058	0.018		0.035	0.058	18,824		23,567	27,588
... 12th Grade or Veteran	0.049		0.067	0.091	0.061		0.067	0.073	25,473		25,532	25,590
... HS Grad	0.020		0.040	0.058	0.030		0.040	0.049	18,260		21,318	24,992
... HS Grad or Non-Veteran	0.026		0.047	0.051	0.023		0.037	0.052	22,687		32,814	42,521
... Other	0.015		0.030	0.053	0.031		0.046	0.061	18,244		24,774	43,213
... None	0.005		0.018	0.059	0.007		0.026	0.052	19,802		40,272	149,084
Years Continuous in USA												
... Used YRSUSA	0.008		0.034	0.058	0.025		0.034	0.063	15,424		22,912	25,529
... No YRSUSA	0.018		0.031	0.057	0.015		0.031	0.057	19,489		25,868	52,379




Figure 12: Task 2 Effect Size and Standard Error